



# Scalable realtime data processing using Kafka Streams

Trayan Iliev

IPT – IT Education Evolved, <http://iproduct.org/>

Copyright © 2003-2023 IPT – Intellectual Products & Technologies Ltd. All rights reserved. Licensed under [CC BY-NC-SA 4.0](https://creativecommons.org/licenses/by-nc-sa/4.0/)



Следете актуалните обяви за Java

DEV.BG



# Machine Learning + Big Data in Real Time + Cloud Technologies => The Future of Intelligent Systems





# Trayan Iliev

- CEO of [IPT – Intellectual Products & Technologies](#) – IT Education Evolved
- Oracle® certified programmer 15+ Y
- 12+ years IT trainer
- End-to-end reactive full-stack apps with Go, Python, Java, Kotlin, TypeScript, React, React Native, Angular, and Vue.js
- Machine Learning (ML) and Deep Learning (DL) with Python, Neuro-Symbolic learning, big data/event processing in realtime (Spark, Kafka, etc.), IoT, уеб и мобилни технологии.
- Voxxed Days, jPrime, Java2Days, jProfessionals, BGOUG, BGJUG, DEV.BG speaker
- Lecturer @ Sofia University – courses: Fullstack React, Angular & TypeScript, Spring 5 Reactive, Distributed Machine Learning, Practical Robotics & IoT
- Robotics / smart-things/ IoT enthusiast, RoboLearn hackathons organizer



Следете актуалните обяви за Java

DEV.BG



# Batch Processing



Extract

Transform

Load



Extract

Load

Transform

Transform

Transform



Следете актуалните обяви за Java

DEV.BG



# Data / Event / Message Streams

“Conceptually, a stream is a (potentially never-ending) **flow of data records**, and a transformation is an operation that takes one or more streams as input, and produces one or more output streams as a result.”

*Apache Flink: Dataflow Programming Model*



Следете актуалните обяви за **Java**

**DEV.BG**



# Data Stream Programming

The idea of **abstracting logic from execution** is hardly new -- it was the dream of **SOA**. And the recent emergence of microservices and containers shows that the dream still lives on.

For developers, the question is whether they want to learn yet **one more layer of abstraction** to their coding. On one hand, there's the elusive promise of a **common API to streaming engines** that in theory should let you mix and match, or swap in and swap out.

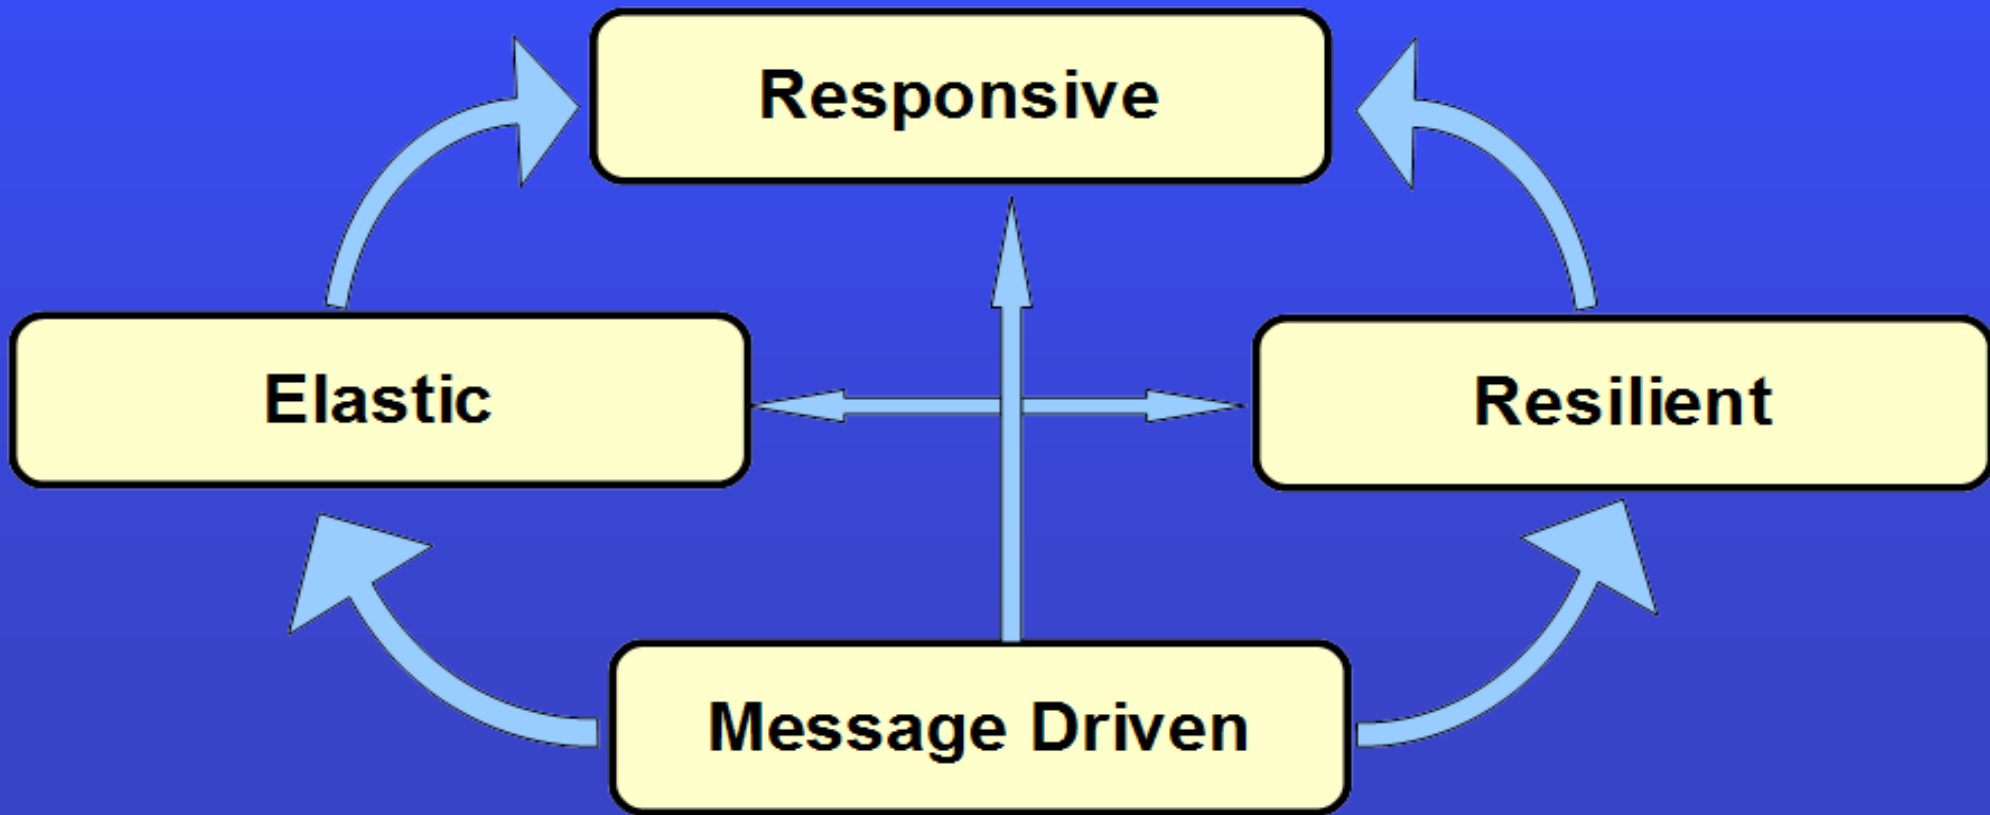
*Tony Baer (Ovum) @ ZDNet - Apache Beam and Spark: New competition for squashing the Lambda Architecture?*





# Reactive Manifesto

<http://www.reactivemanifesto.org>



Следете актуалните обяви за Java

DEV.BG



# Scalable, Massively Concurrent

- **Message Driven** – asynchronous message-passing allows to establish a boundary between components that ensures loose coupling, isolation, location transparency, and provides the means to delegate errors as messages [Reactive Manifesto].
- The main idea is to separate concurrent producer and consumer workers by using **message queues**.
- Message queues can be **unbounded** or **bounded** (limited max number of messages)
- **Unbounded** message queues can present memory allocation problem in case the producers outrun the consumers for a long period → **OutOfMemoryError**









# Synchronous vs. Asynchronous IO

Synchronous



Asynchronous






Следете актуалните обяви за Java

DEV.BG

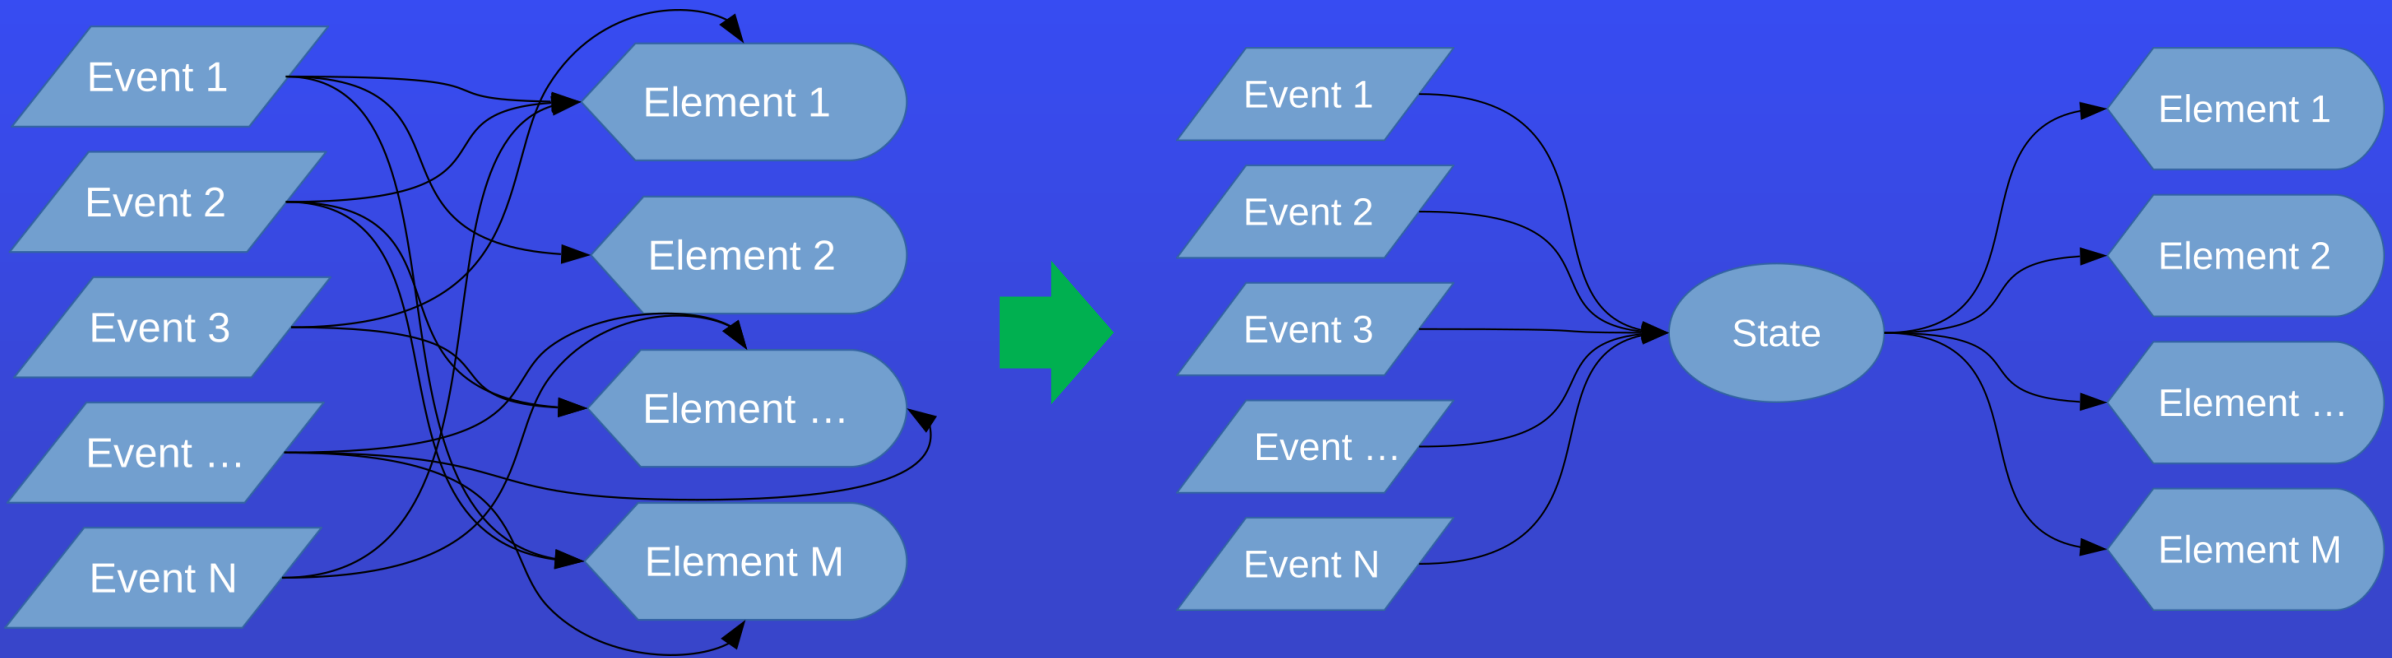


# Stream Topology => Direct Acyclic Graph





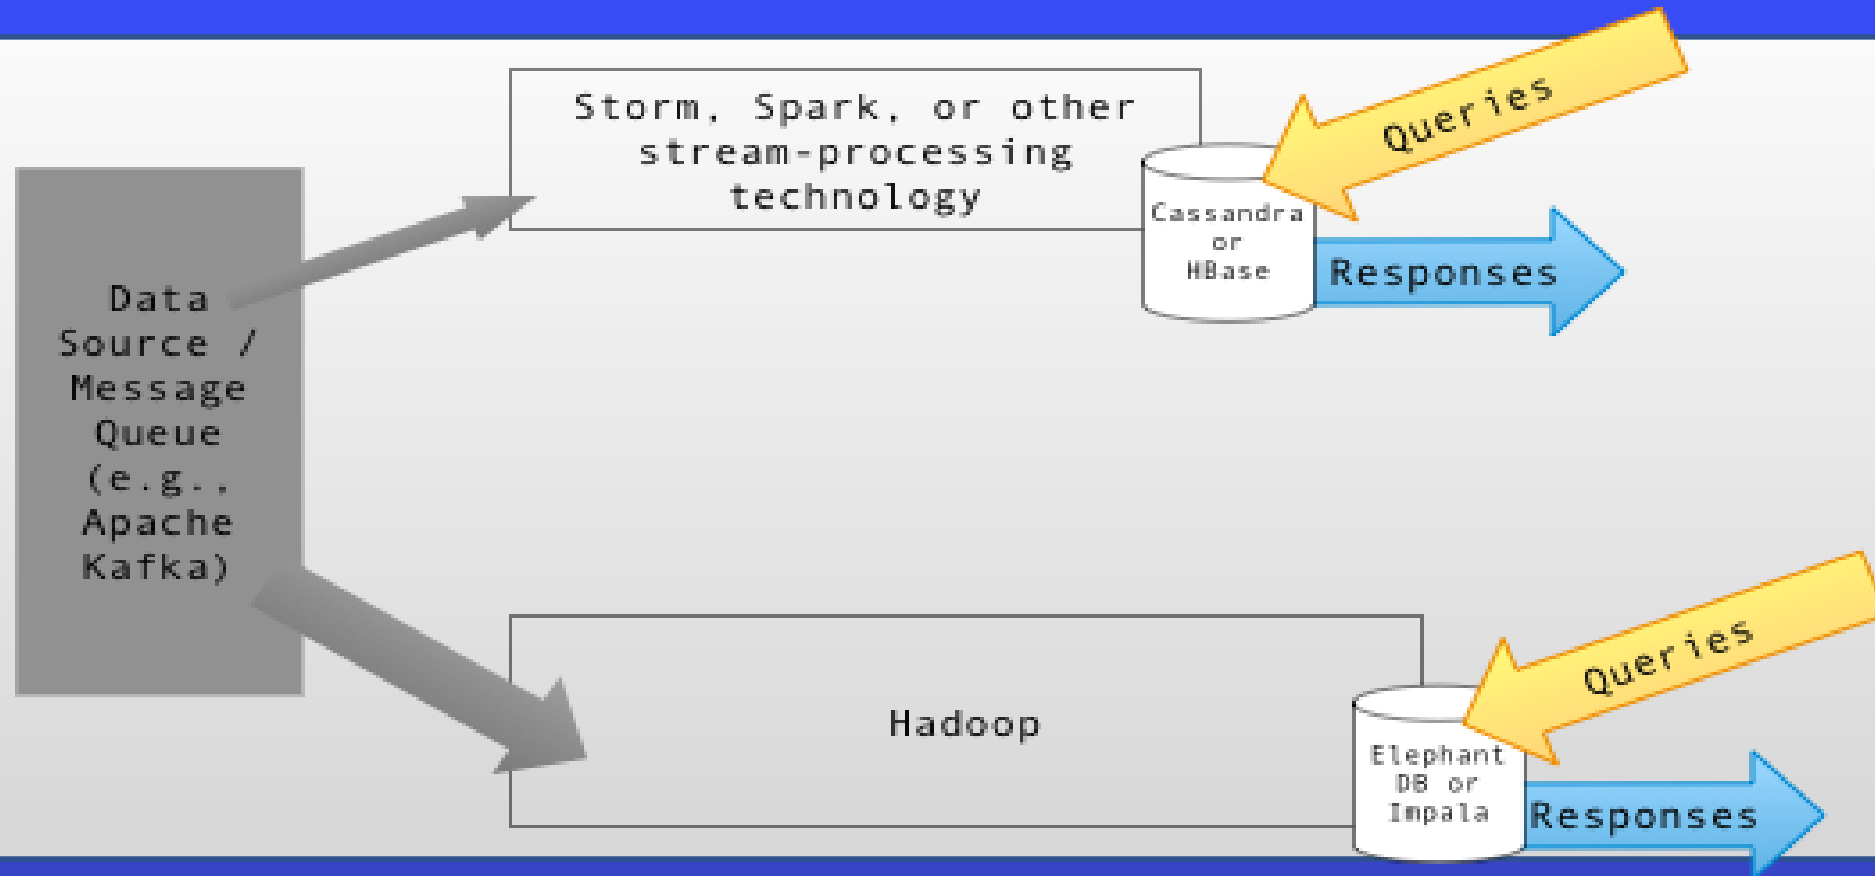
# Event Sourcing – Events vs. State (Snapshots)





# Lambda Architecture - I

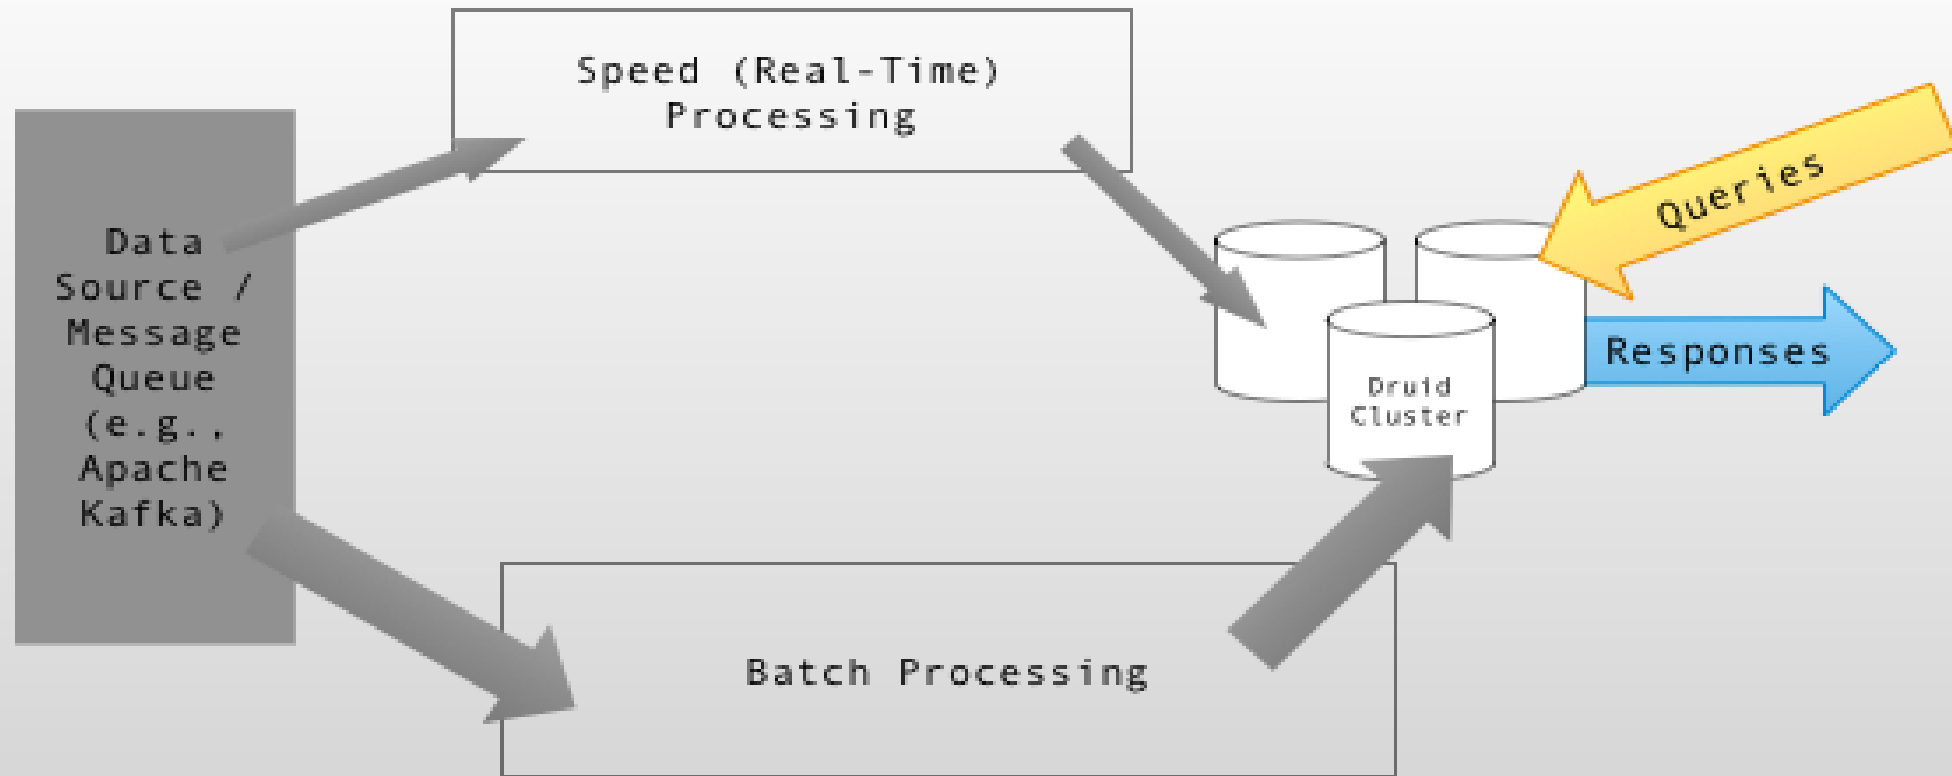
Query =  $\lambda$  (Complete data) =  $\lambda$  (live streaming data) \*  $\lambda$  (Stored data)



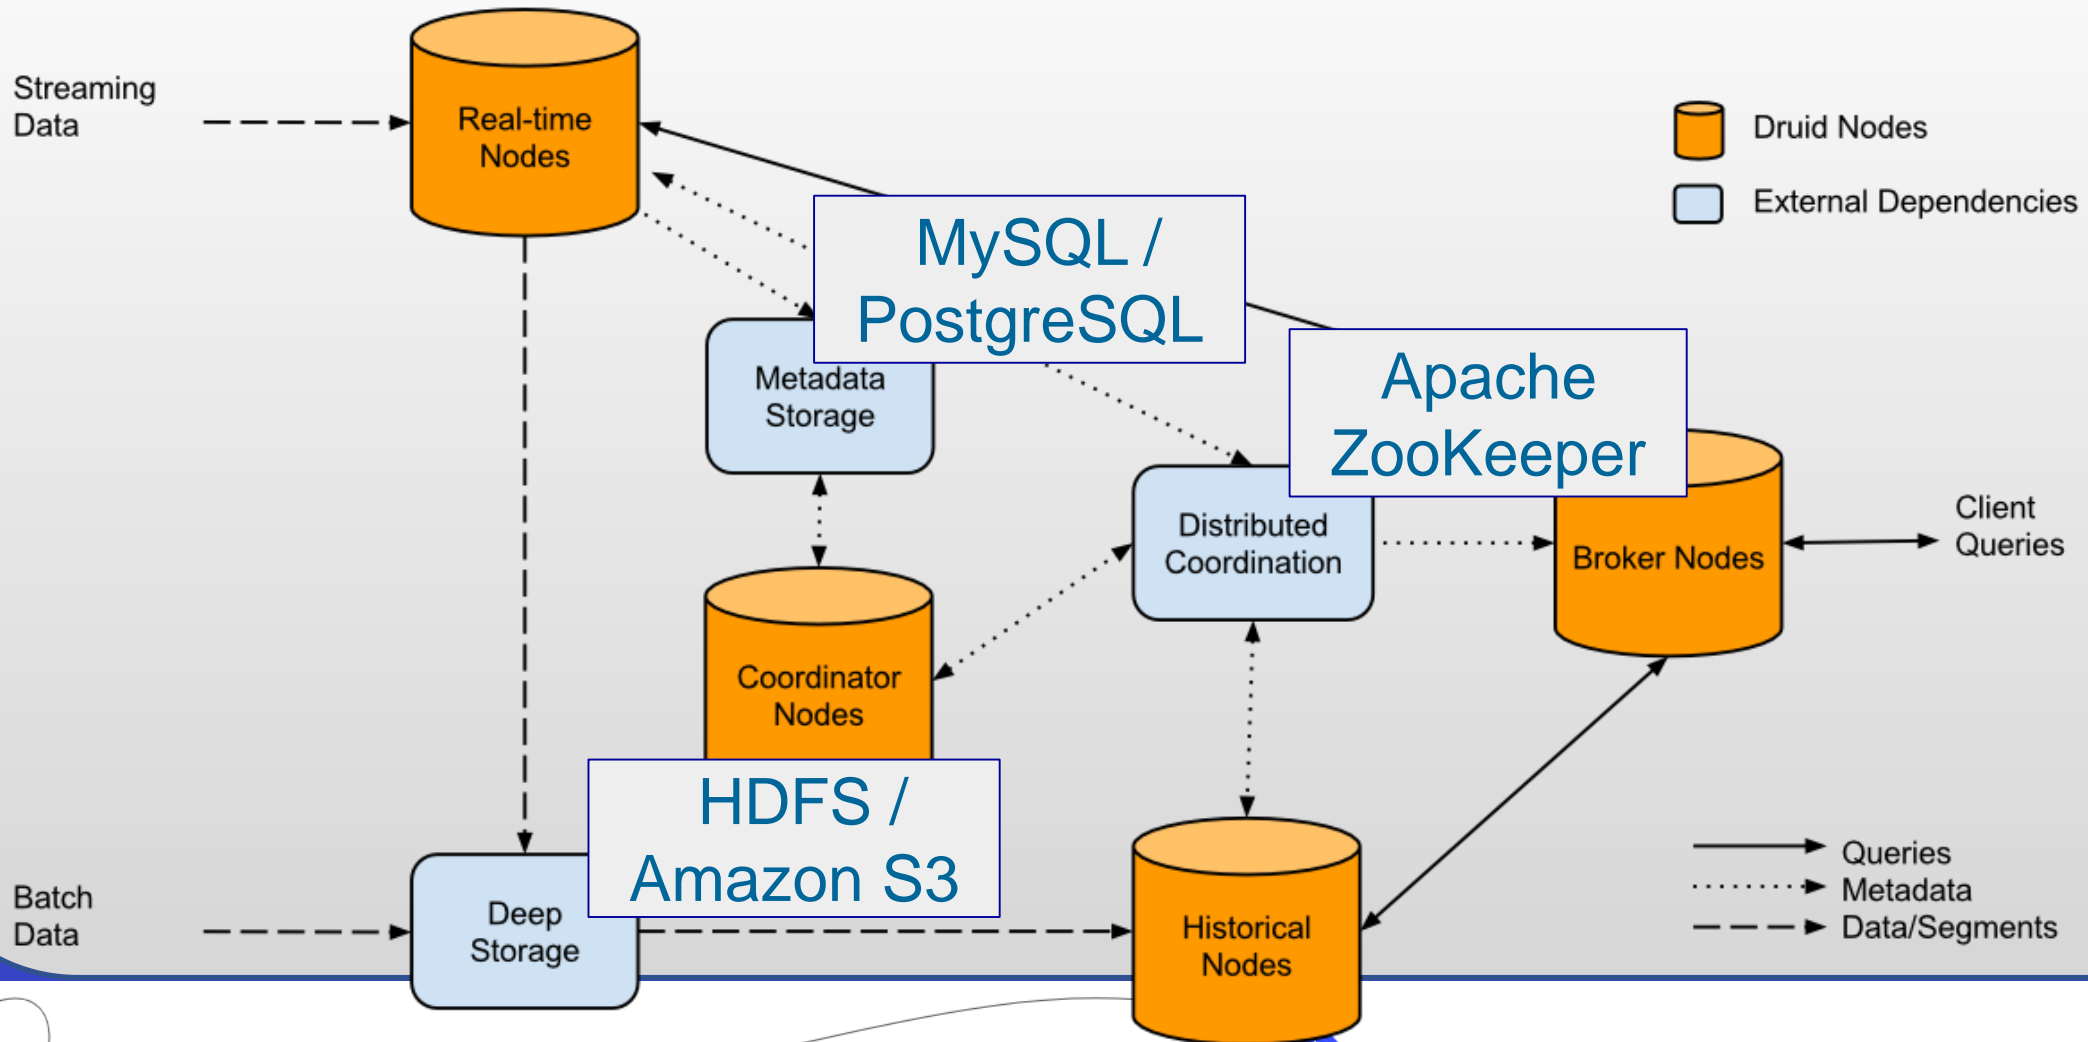


# Lambda Architecture - II

Query =  $\lambda$  (Complete data) =  $\lambda$  (live streaming data) \*  $\lambda$  (Stored data)



# Lambda Architecture - Druid Distributed Data Store





# Kappa Architecture - I

**Query = K (New Data) = K (Live streaming data)**

- Proposed by Jay Kreps in 2014
- Real-time processing of distinct events
- Drawbacks of Lambda architecture:
  - It can result in coding overhead due to comprehensive processing
  - Re-processes every batch cycle which may not be always beneficial
  - Lambda architecture modeled data can be difficult to migrate
- Canonical data store in a Kappa Architecture system is an **append-only immutable log** (like **Kafka, Pulsar**)


$\lambda$  vs  $\kappa$





# Kappa Architecture -II

- Multiple **data events** or queries are **logged in a queue** to be catered against a **distributed file system storage** or history.
- The order of the events and queries is not predetermined. Stream processing platforms can **interact with database at any time**.
- It is **resilient** and **highly available** as handling **terabytes of storage** is required for each node of the system to support replication.
- **Machine learning** is done on the **real time basis**





# Zeta Architecture

Main characteristics of Zeta architecture:

- file system (HDFS, S3, GoogleFS),
- realtime data storage (HBase, Spanner, BigTable),
- modular processing model and platform (MapReduce, Spark, Drill, BigQuery),
- containerization and deployment (cgroups, Docker, Kubernetes, etc.),
- Software solution architecture (serverless computing – e.g. Amazon Lambda)
- Recommender systems and machine learning




# Distributed Stream Processing – Apache Projects:

- **Apache Spark** is an open-source cluster-computing framework. **Spark Streaming**, **Spark Mllib**, **Spark ML** (ML pipelines using Dataframes)
- **Apache Storm** is a distributed stream processing – streams DAG
- **Apache Samza** is a distributed real-time stream processing framework.




# Distributed Stream Processing – Apache Projects:

- **Apache Flink** - open source stream processing framework – stateful computations over data streams - **Flink ML**: Machine Learning library
- **Apache Kafka** - open-source stream processing (**Kafka Streams**), real-time, low-latency, high-throughput, massively scalable pub/sub
- **Apache Beam** – unified batch and streaming, portable, extensible











# Example: Internet of Things (IoT) & Robotics


3D accelerometers, gyros,  
and compass MinIMU-9 v2

USB Stereo  
Speakers - 5V

LiPo Powebank  
15000 mAh

Pololu DRV8835  
Dual Motor Driver  
for Raspberry Pi


Arduino Leonardo clone  
A-Star 32U4 Micro



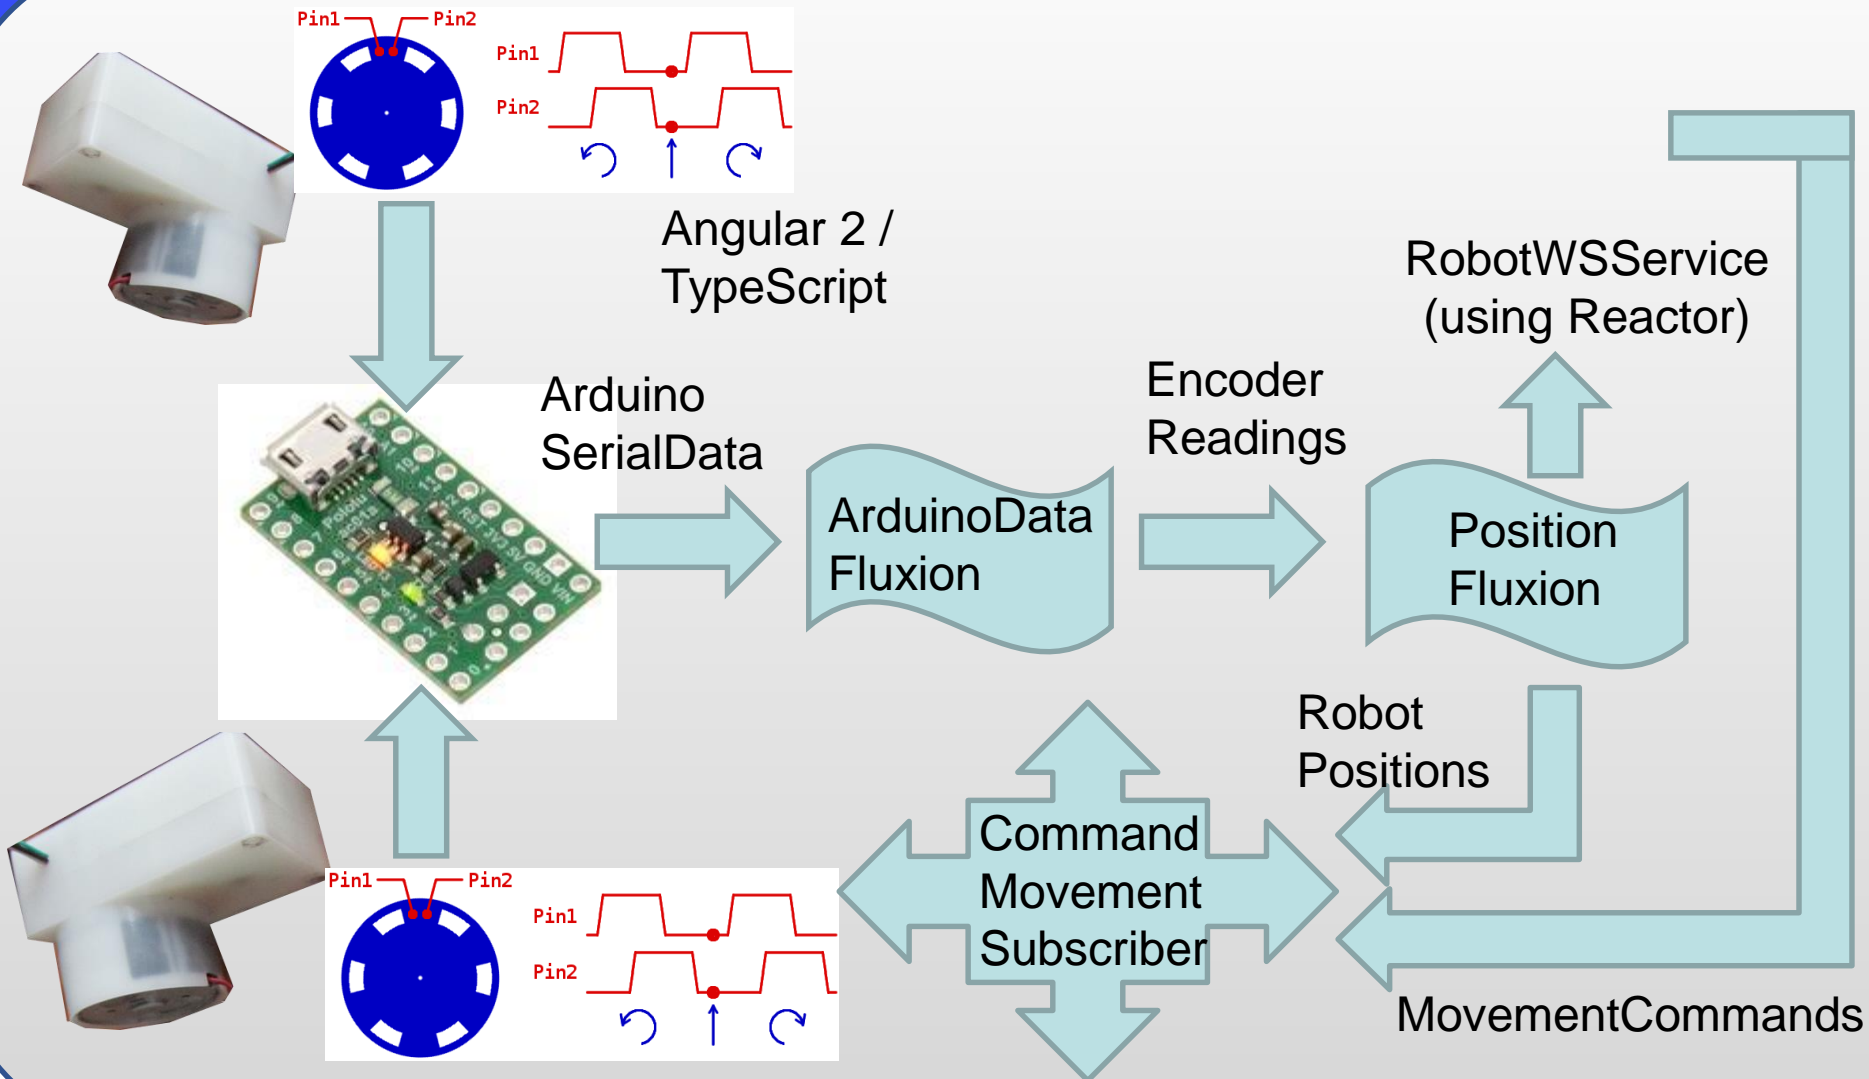
Следете актуалните обяви за Java

DEV.BG





# Adaptive Robot Control using Reactive Streams







# Apache Kafka Distributed Streaming Platform

- Kafka achieves **high-throughput, low-latency, durability, and near-limitless scalability** by maintaining a distributed system based on commit logs, delegating key responsibility to clients, optimizing for **batches** and allowing for **multiple concurrent consumers** per message.
- **Publish and subscribe (Pub/Sub)** to streams of records – similar to a message queue or enterprise messaging system
- **Store streams of records** – in a fault-tolerant and durable way
- **Process streams of records** – as they occur (in real-time)





# Two Types of Applications for Kafka

- Building real-time streaming data pipelines that **reliably get data between systems or applications**
- Building real-time **streaming applications** that **transform or react to the streams of data** – Kafka Streams





# Apache Kafka Typical Use-Cases

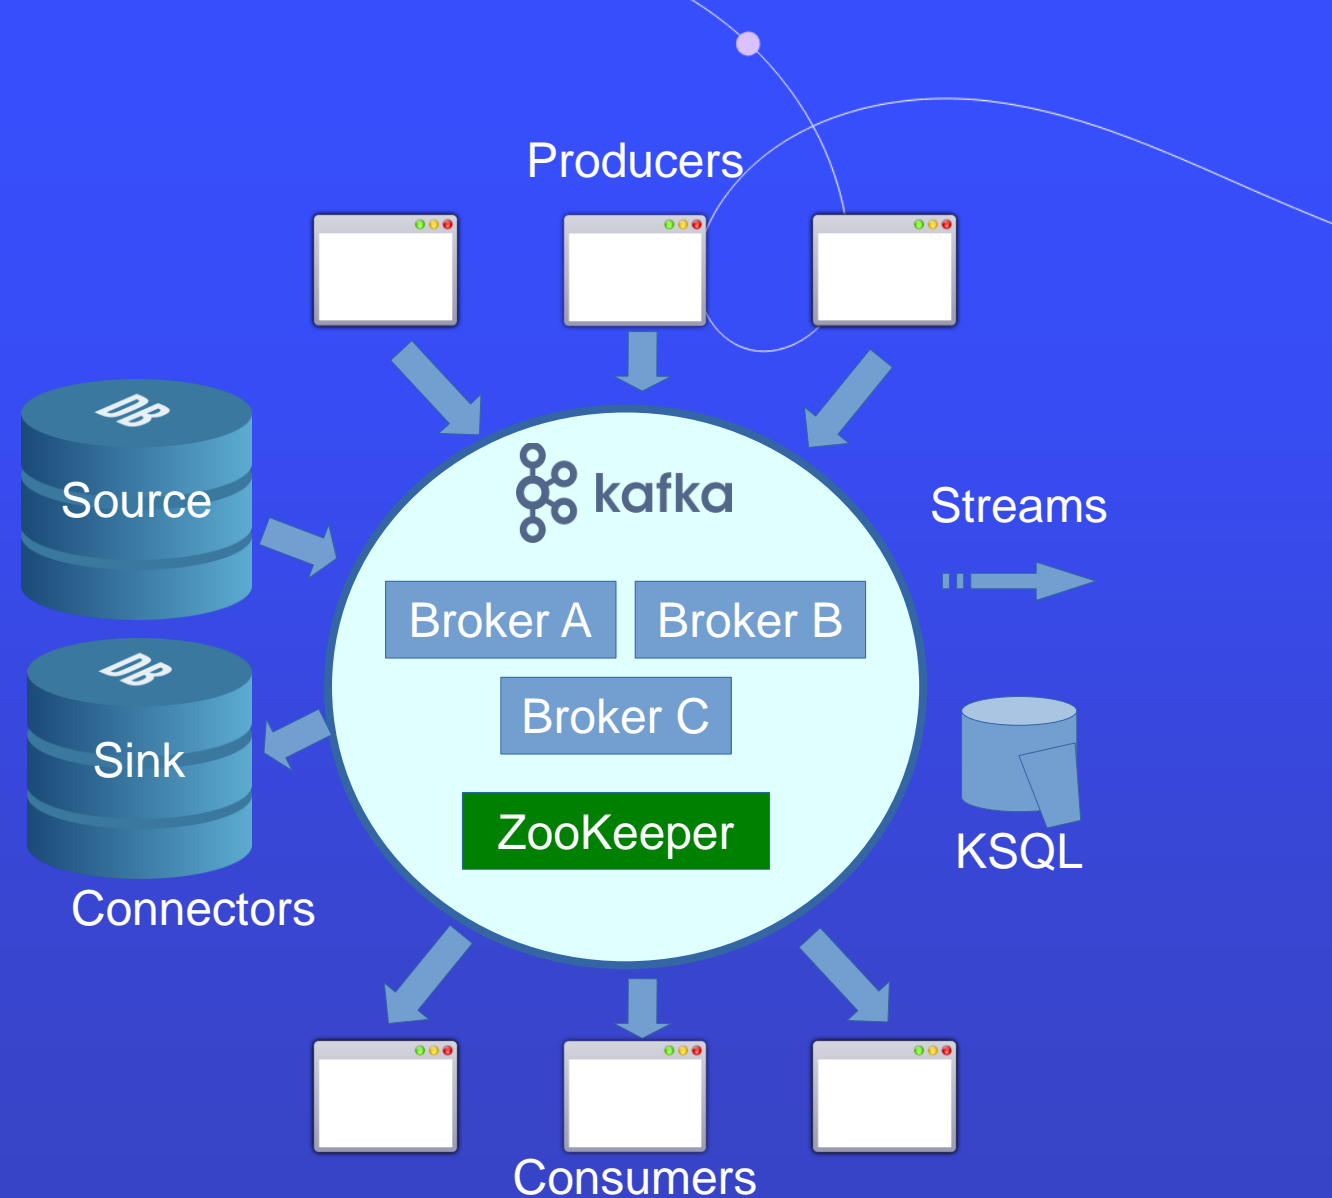
- IoT, telemetry, and sensor networks
- Positional data / Logistics - supply chain and transportation alerts
- Service/process monitoring - aggregating metrics and logs from distributed servers and applications (Event-driven SOA)
- Real-time analytics, fraud detection – processing of business/customer events in real time
- Click stream analytics, real-time predictive analytics
- Stock-trading analysis





# Kafka Main Concepts

- Kafka is run as a **cluster** on one or more servers (**brokers**) that can span multiple datacenters.
- The Kafka cluster stores **streams of records** in categories called **topics**.
- Each **record** consists of a **key**, **value**, and **timestamp**.





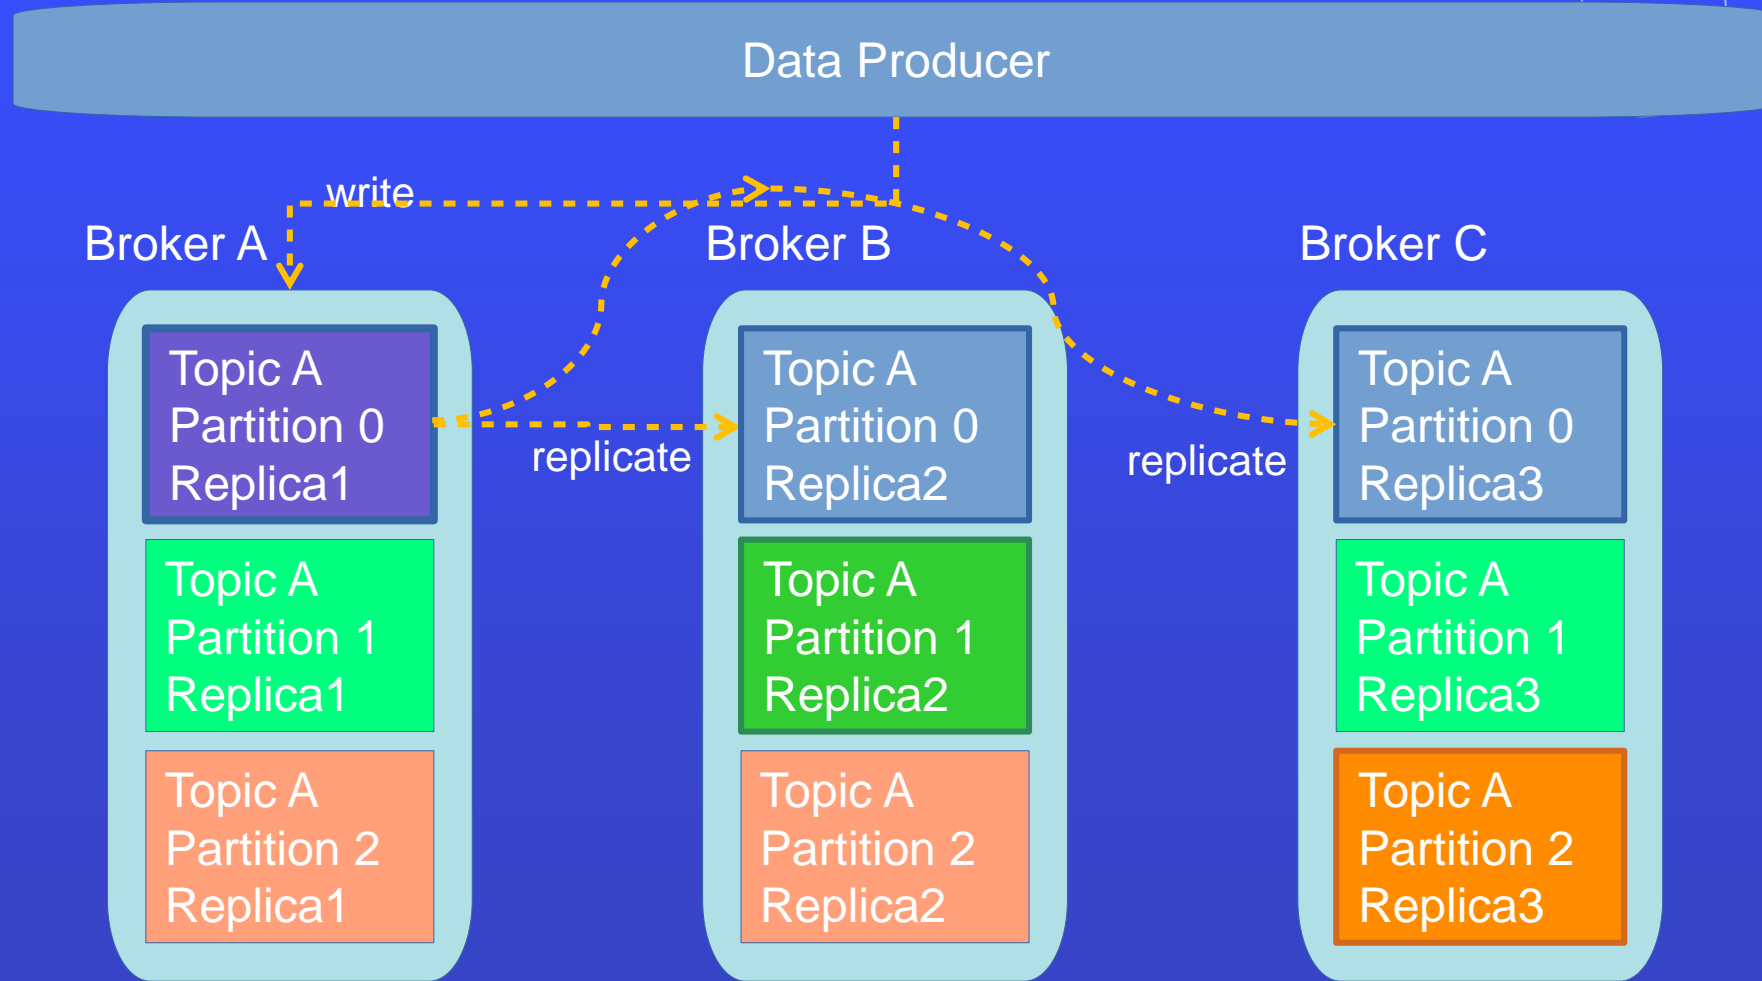
# Kafka Core APIs

- **Producer API** - publish a stream of records to one or more Kafka topics.
- **Consumer API** - subscribe to one or more topics and process the stream of records produced to them.
- **Streams API** - a stream processor, consuming an input stream from one or more topics and producing an output stream to one or more output topics, effectively transforming the input streams to output streams.
- **Connector API** allows building and running reusable producers or consumers that connect Kafka topics to existing applications or data systems – e.g. connector to a DB might capture every change in a table



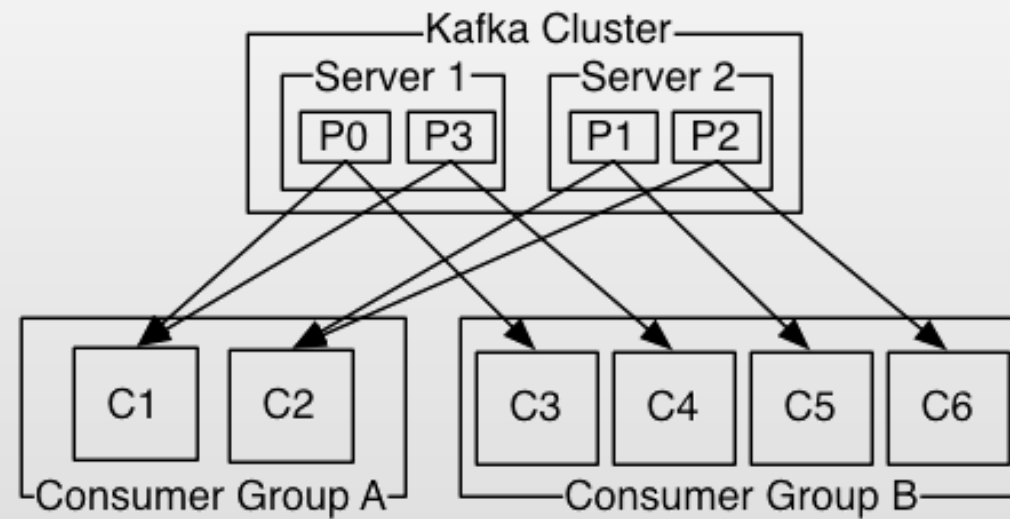
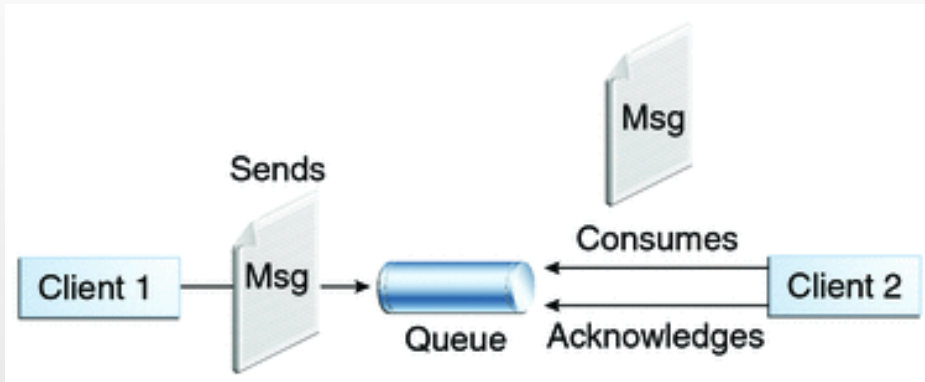


# Kafka Data Replication





# Kafka as a Messaging System





# Kafka Streams

- By combining **storage** and **low-latency subscriptions**, streaming applications can treat both **past and future data** the same way. That is a single application can process historical, stored data but rather than ending when it reaches the last record it can **keep processing as future data arrives**. This is a generalized notion of stream processing that subsumes batch processing as well as message-driven applications => **Kappa architecture**
- Likewise for **streaming data pipelines** the combination of subscription to real-time events make it possible to use Kafka for **very low-latency pipelines**; but the ability to store data reliably make it possible to use it for critical data where the **delivery of data must be guaranteed** or for integration with offline systems that load data only periodically or may go down for extended periods of time for maintenance.







# Why you'll love using Kafka Streams?

- Elastic, highly scalable, fault-tolerant
- Deploy to containers, VMs, bare metal, cloud
- Equally viable for small, medium, & large use cases
- Fully integrated with Kafka security
- Write standard Java and Scala applications
- Exactly-once processing semantics
- No separate processing cluster required
- Develop on Mac, Linux, Windows





# Kafka Streams Advantages - I

- Designed as a **simple** and **lightweight** client library, which can be easily embedded in any Java application and integrated with any existing packaging, deployment and operational tools that users have for their streaming applications.
- Has **no external dependencies** on systems other than Apache Kafka itself as the internal messaging layer; notably, it uses Kafka's partitioning model **to horizontally scale processing** while maintaining **strong ordering guarantees**.
- Supports **fault-tolerant local state**, which enables very **fast and efficient** stateful operations like **windowed joins** and **aggregations**.





# Kafka Streams Advantages - II

- Supports **exactly-once processing semantics** to guarantee that each record will be processed once and only once even when there is a failure on either Streams clients or Kafka brokers in the middle of processing.
- Employs **one-record-at-a-time** processing to achieve **millisecond processing latency**, and supports event-time based **windowing operations** with **out-of-order arrival** of records.
- Offers necessary stream processing primitives, along with a high-level **Streams DSL** and a **low-level Processor API**.





# Stream Processing Topology - I

- A **stream** is the most important abstraction provided by Kafka Streams: it represents an unbounded, continuously updating data set. A stream is an ordered, replayable, and fault-tolerant sequence of **immutable data** records, where a **data record** is defined as a **key-value pair**.
- A **stream processing application** is any program that makes use of the **Kafka Streams** library. It defines its computational logic through one or more processor topologies, where a **processor topology** is a **graph of stream processors (nodes)** that are connected by streams (edges).
- A **stream processor** is a **node in the processor topology**; it represents a processing step to transform data in streams by receiving **one input record at a time** from its **upstream processors** in the topology, **applying its operation** to it, and may subsequently produce one or more output records to its **downstream processors**.





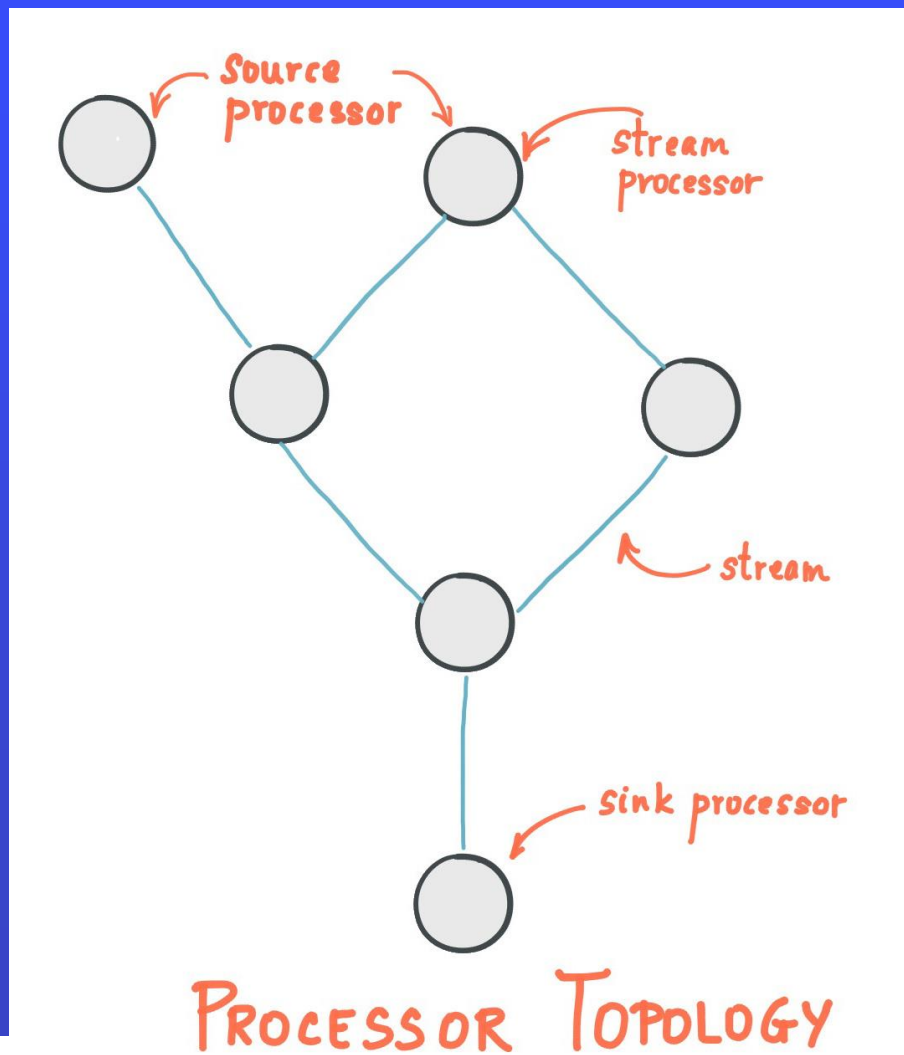
# Types of Processors

- **Source Processor:** A source processor is a special type of stream processor that does not have any upstream processors. It produces an input stream to its topology from one or multiple Kafka topics by consuming records from these topics and forwarding them to its down-stream processors.
- **Sink Processor:** A sink processor is a special type of stream processor that does not have down-stream processors. It sends any received records from its up-stream processors to a specified Kafka topic.
- Note that in normal processor nodes other remote systems can also be accessed while processing the current record. Therefore the processed results can either be streamed back into Kafka or written to an external system.





# Kafka Stream Processing - DAG





# Time in Kafka Streams

- A critical aspect in stream processing is the **notion of time**, and how it is modeled and integrated. For example, some operations such as **windowing** are defined based on **time boundaries**:
- **Event time** - the point in time when an event or data record occurred, i.e. was originally created "at the source".
- **Processing time** - the point in time when the event or data record happens to be processed by the stream processing application, i.e. when the record is being consumed.
- **Ingestion time** – time when an event or data record is stored in a topic partition by a Kafka broker. The difference to event time is that this ingestion timestamp is generated when the record is appended to the target topic by Kafka broker, not when the record is created at source.





# Time in Kafka Streams: Configuration

- **log.message.timestamp.type** – define whether the timestamp in the message is message create time or log append time. The value should be either `CreateTime` or `LogAppendTime`
- **log.message.timestamp.difference.max.ms** – The maximum difference allowed between the timestamp when a broker receives a message and the timestamp specified in the message. If `log.message.timestamp.type=CreateTime`, a message will be rejected if the difference in timestamp exceeds this threshold. This configuration is ignored if `log.message.timestamp.type=LogAppendTime`. The maximum timestamp difference allowed should be no greater than `log.retention.ms` to avoid unnecessarily frequent log rolling.







# Custom TimestampExtractor

@Slf4j

```
public class CustomTimeExtractor implements TimestampExtractor {
```

@Override

```
public long extract(ConsumerRecord<Object, Object> record, long partitionTime) {
```

```
    final long timestamp = record.timestamp();
```

*// `TemperatureReading` is your own custom class, which we assume has a method that returns  
// the embedded timestamp (in milliseconds).*

```
    var myReading = (TemperatureReading) record.value();
```

```
    if (myReading != null) {
```

```
        return java.sql.Timestamp.valueOf(myReading.getTimestamp()).getTime();
```

```
    }
```

```
    else {
```

*// Kafka allows `null` as message value. How to handle such message values  
// depends on your use case. In this example, we decide to fallback to  
// wall-clock time (= processing-time).*

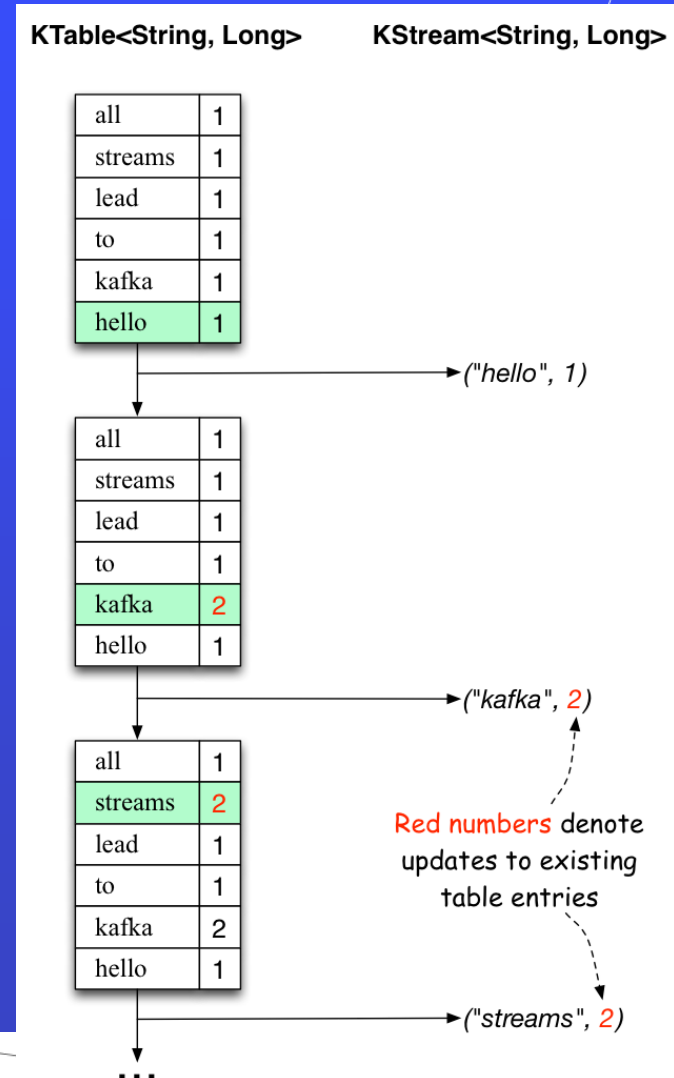
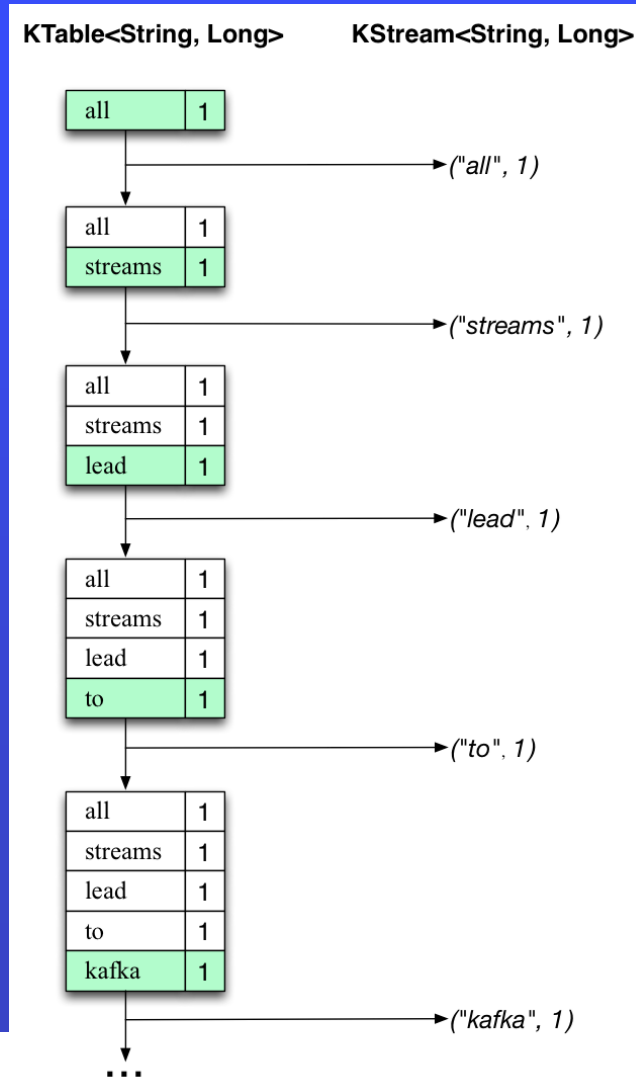
```
        return System.currentTimeMillis();
```

```
    }
```

```
}
```



# Kafka Stream Processing Example





# Kafka Streams Dependencies

```
dependencies {  
    implementation 'org.apache.kafka:kafka-clients:3.4.0'  
    implementation 'org.apache.kafka:kafka-streams:3.4.0'  
    ...  
}
```





# Kafka Streams Code Skeleton

```
public static void main(String[] args) {  
    // Use the builders to define the actual processing topology, e.g. to specify from which input topics to  
    read,  
    // which stream operations (filter, map, etc.) should be called, and so on.  
  
    StreamsBuilder builder = ...; // when using the DSL  
    Topology topology = builder.build();  
    //  
    // OR  
    //  
    Topology topology = ...; // when using the Processor API  
  
    // Use the configuration to tell your application where the Kafka cluster is,  
    // which Serializers/Deserializers to use by default, to specify security settings, and so on.  
    Properties props = ...;  
    KafkaStreams streams = new KafkaStreams(topology, props);  
    // Add shutdown hook to stop the Kafka Streams threads. You can optionally provide a timeout to  
    `close`.  
    Runtime.getRuntime().addShutdownHook(new Thread(streams::close));  
}
```



# Stream Partitions and Tasks

- Kafka **messaging layer partitions** data for storing and transporting it.
- **Kafka Streams partitions** data for processing it.
- In both cases, this partitioning is what enables data **locality, elasticity, scalability, high performance, and fault tolerance**. Kafka Streams uses the concepts of **partitions** and **tasks** as **logical units of its parallelism** model based on **Kafka topic partitions**.
- Each **stream partition** is a totally ordered sequence of data records and maps to a **Kafka topic partition**.
- A **data record** in the stream maps to a Kafka **message** from that topic.
- The **keys** of data records determine the **partitioning of data** in both **Kafka** and **Kafka Streams** -how data is routed to specific topic partitions.





# Stream Partitions and Tasks - II

- An application's **processor topology** is scaled by breaking it into multiple **tasks**.
- Kafka Streams creates a **fixed number of tasks** based on the **input stream partitions** for the application, with **each task assigned a list of partitions** from the input streams (i.e., Kafka topics).
- The **assignment of partitions to tasks never changes** so that each task is a **fixed unit of parallelism** of the application.
- Tasks can then instantiate their own **processor topology** based on the assigned partitions; they also maintain a **buffer for each of its assigned partitions** and **process messages one-at-a-time** from these record buffers.
- As a result stream tasks can be processed independently and in parallel without manual intervention.





# Stream Partitions and Tasks - III

- Kafka Streams is **NOT a resource manager**, but a **library** that "runs" anywhere its stream processing application runs.
- **Multiple instances of the application** are executed either **on the same machine**, or spread across **multiple machines** and tasks can be distributed automatically by the library to those running application instances.
- **Assignment of partitions to tasks never changes** - if an application instance fails, all its assigned tasks will be automatically restarted on other instances and continue to consume from the same stream partitions.
- **Topic partitions are assigned to tasks**, and **tasks are assigned to all threads over all instances**, in a best-effort attempt to trade off load-balancing and stickiness of stateful tasks. For this assignment, Kafka Streams uses the **StreamsPartitionAssignor** class.





# StreamsPartitionAssignor Tasks Assignment


1. Decode the subscriptions to assemble the metadata for each client and check for version probing.
2. Check all repartition source topics and use internal topic manager to make sure they have been created with the right number of partitions. Also verify and/or create any changelog topics with the correct number of partitions.
3. Use the partition grouper to generate tasks along with their assigned partitions, then use the configured **TaskAssignor** to construct the mapping of tasks to clients.
4. Construct the global mapping of host to partitions to enable query routing.
5. Within each client, assign tasks to consumer clients.








# Kafka Streams Partitions and Tasks - I





# Kafka Streams Partitions and Tasks - II





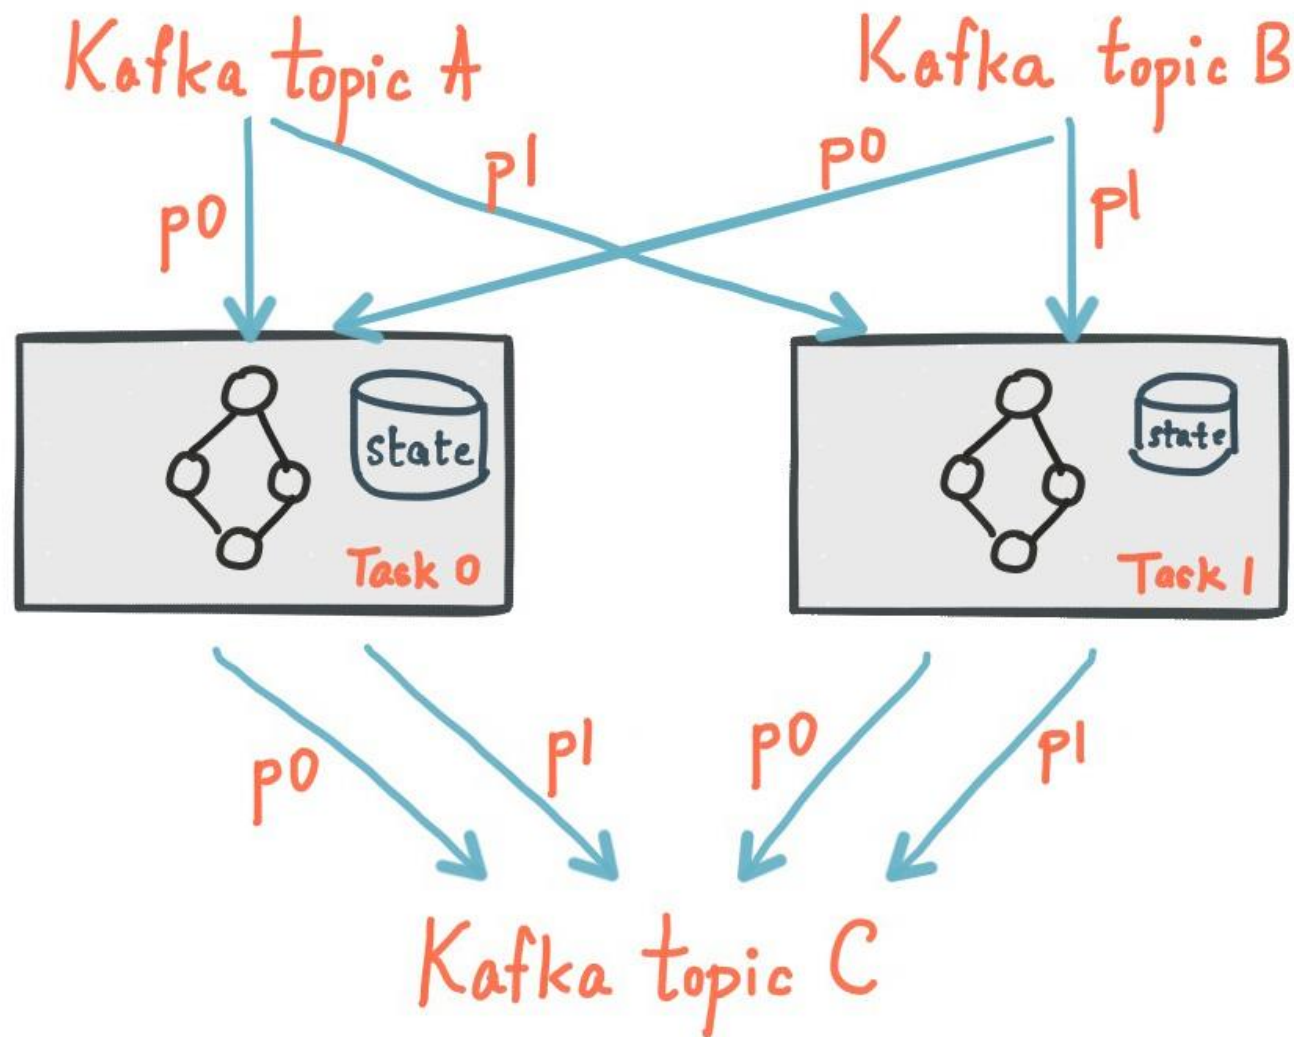
# Tasks Threading Model

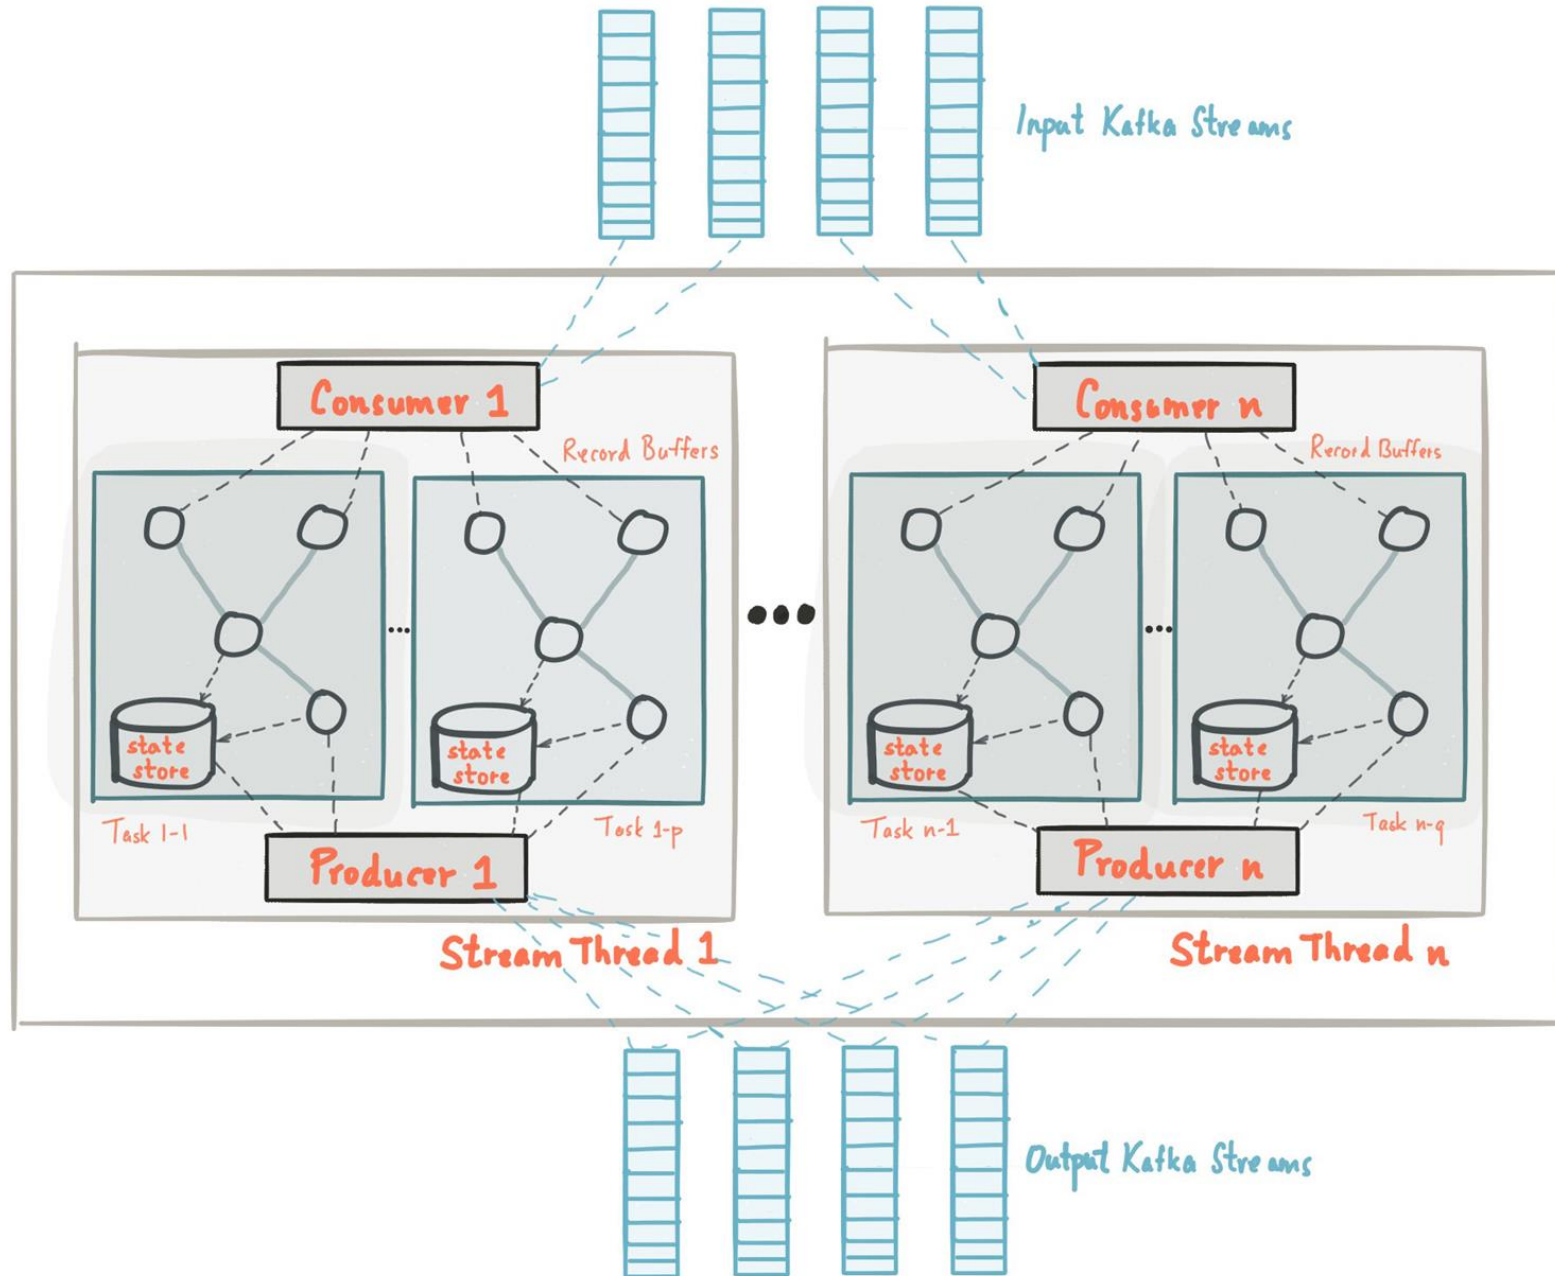
- Starting more **stream threads** or more **instances of the application** merely amounts to **replicating the topology** and having it process a **different subset of Kafka partitions**, effectively parallelizing processing.
- It is worth noting that there is **no shared state amongst the threads**, so **no inter-thread coordination is necessary**.
- This makes it very simple to **run topologies in parallel** across the **application instances and threads**.
- The **assignment of Kafka topic partitions** amongst the various stream threads is transparently handled by **Kafka Streams** + Kafka coordination.
- You can start **as many threads of the application** as there **are input topic partitions** so that, across all running instances of an application, every thread (or rather, the tasks it runs) has at least one input partition to process.





# Kafka Streams Partitions and Tasks - III









# Kafka Streams DSL & Processor API

- **Processor API** - allows developers to define and connect custom processors and to interact with state stores. With the Processor API, you can define arbitrary stream processors that process one received record at a time, and connect these processors with their associated state stores to compose the processor topology that represents a customized processing logic.
- **Processor API** can be used to implement both stateless as well as stateful operations, where the latter is achieved through the use of state stores.
- **Kafka Streams DSL** (Domain Specific Language) is built on top of the Streams Processor API. It is the recommended for most users, especially beginners. Most data processing operations can be expressed in just a few lines of DSL code.
- Combining the **DSL** and the **Processor API** – you can combine the convenience of the **DSL** with the power and flexibility of the **Processor API** as described in [Applying processors and transformers \(Processor API integration\)](#).





# Kafka Streams DSL: KStreams

- Only the **Kafka Streams DSL** has the notion of a **KStream**.
- **KStream** is an abstraction of a record stream, where each data record represents a self-contained datum in the unbounded data set. Using the table analogy, data records in a record stream are always interpreted as an **"INSERT"** -- thing: adding more entries to an append-only ledger -- because no record replaces an existing row with the same key. Examples are a credit card transaction, a page view event, or a server log entry.
- To illustrate, let's imagine the following two data records are being sent to the stream:  
**("alice", 1) --> ("alice", 3)**
- If your stream processing application were to sum the values per user, it would return **4** for **alice**. Why? Because the second data record would not be considered an update of the previous record. Compare this behavior of **KStream** to **KTable** in next slide, which would return **3** for **alice**.





# Kafka Streams DSL: KTables

- Only the **Kafka Streams DSL** has the notion of a **KTable**.
- **KTable** is an abstraction of a changelog stream, where each data record represents an update. More precisely, the value in a data record is interpreted as an **"UPDATE"** of the last value for the same record key, if any (if a corresponding key doesn't exist yet, the update will be considered an **INSERT**). Using the table analogy, a data record in a changelog stream is interpreted as an **UPSERT** aka **INSERT/UPDATE** because any **existing row with the same key is overwritten**. Also, **null values (tombstones)** are interpreted in a special way: a **record with a null value** represents a **"DELETE"** or **tombstone** for the record's key.
- To illustrate, let's imagine the following two data records are being sent to the stream: **("alice", 1) --> ("alice", 3)**
- If a stream processing application is summing the values per user, it would return **3** for **alice**. Why? Second record would be considered update previous.







# KTables and Log Compaction

- Another way of thinking about **KStream** and **KTable** is as follows: If you were to store a **KTable** into a Kafka topic, you'd probably want to enable Kafka's **log compaction** feature, e.g. to save storage space.
- However, it would not be safe to enable log compaction in the case of a **KStream** because, as soon as log compaction would begin purging older data records of the same key, it would break the semantics of the data. E.g. you'd suddenly get a **3** for **alice** instead of a **4** because of log compaction. Hence log compaction is perfectly safe for a **KTable** (**changelog stream**) but it is a mistake for a **KStream** (**record stream**).
- Example: **Change Data Capture (CDC)** records in the changelog of a relational DB, representing which row in database table was inserted/ updated/ deleted.
- **KTable** also provides an ability to **look up current values** of data records **by keys**. Table-lookup is available through **join operations** & **Interactive Queries**.





# Kafka Streams DSL: GlobalKTable

- **GlobalKTable** is an abstraction of a **changelog stream**, where each data record represents an update.
- **GlobalKTable** differs from a **KTable** in the data that they are being populated with, i.e. which data from the underlying Kafka topic is being read into the respective table. Slightly simplified, imagine you have an **input topic** with **5 partitions**. In your application, you want to read this topic into a table. You want to run your application across **5 application instances** for **maximum parallelism**.
- If input topic read into a **KTable**, then "local" **KTable** instance of each application instance will be populated with data from only 1 partition of the topic 5 partitions.
- If input topic read into a **GlobalKTable**, then the local **GlobalKTable** instance of each application instance will be populated with **data from all topic**.
- **GlobalKTable** provides the ability to **look up current values** of data records **by keys**. This table-lookup functionality is available through **join operations**. Note that a **GlobalKTable** has **no notion of time** in contrast to a **KTable**.





# Benefits and Downsides of Using GlobalKTable

- Benefits:
  - More convenient and/or efficient **joins**: Notably, global tables allow you to perform **star joins**, they support "foreign-key" **lookups** (i.e., you can lookup data in the table **not just by record key**, but also by data in the **record values**), and they are more efficient when **chaining multiple joins**. Also, when **joining against a global table**, the input data does not need to be **co-partitioned**.
  - Can be used to "**broadcast**" information to all the running instances of your application.
- Downsides of global tables:
  - **Increased local storage consumption** compared to the (partitioned) KTable because the entire topic is tracked.
  - **Increased network and Kafka broker load** compared to the (partitioned) KTable because the entire topic is read.





# Streams DSL: Creating a Stream

```
import org.apache.kafka.common.serialization.Serdes;
import org.apache.kafka.streams.StreamsBuilder;
import org.apache.kafka.streams.kstream.Consumed;
import org.apache.kafka.streams.kstream.KStream;

public class Temp {
    public static void main(String[] args) {
        StreamsBuilder builder = new StreamsBuilder();
        KStream<String, Long> wordCounts = builder.stream(
            "word-counts-input-topic", /* input topic */
            Consumed.with(
                Serdes.String(), /* key serde */
                Serdes.Long() /* value serde */
            ));
    }
}
```





# Streams DSL: Creating GlobalKTable

```
import org.apache.kafka.common.serialization.Serdes;
import org.apache.kafka.common.utils.Bytes;
import org.apache.kafka.streams.StreamsBuilder;
import org.apache.kafka.streams.kstream.GlobalKTable;
import org.apache.kafka.streams.kstream.Materialized;
import org.apache.kafka.streams.state.KeyValueStore;

public class Temp {
    public static void main(String[] args) {
        StreamsBuilder builder = new StreamsBuilder();
        GlobalKTable<String, Long> wordCounts = builder.globalTable(
            "word-counts-input-topic",
            Materialized.<String, Long, KeyValueStore<Bytes, byte[]>>as(
                "word-counts-global-store" /* table/store name */)
                .withKeySerde(Serdes.String()) /* key serde */
                .withValueSerde(Serdes.Long()) /* value serde */
        );
    }
}
```



# Streams DSL KStream and KTable Transformations

- **KStream** is an abstraction of a record stream of **KeyValue** pairs, i.e., each record is an independent entity/event in the real world. For example a user **X** might buy two items **I1** and **I2**, and thus there might be two records **<K:I1>**, **<K:I2>** in the stream.
- A **KStream** is either defined from **one or multiple Kafka topics** that are **consumed message by message**, or the result of a **KStream transformation**.
- A **KTable** can also be converted into a **KStream**.
- A **KStream** can be **transformed** record by record, **joined** with another **KStream**, **KTable**, **GlobalKTable**, or can be **aggregated** into a **KTable**. **Kafka Streams DSL** can be mixed-and-matched with **Processor API (PAPI)** (c.f. Topology) via **process(...)**, **transform(...)**, and **transformValues(...)**.





# Processor API (PAPI) Example - I

```
public class WordCountProcessor implements Processor<String, String, String, String> {  
    private KeyValueStore<String, Long> kvStore;  
    private ProcessorContext<String, String> context;  
  
    @Override  
    public void init(ProcessorContext<String, String> context) {  
        this.context = context;  
        kvStore = context.getStateStore("inmemory-word-counts");  
    }  
  
    @Override  
    public void close() {  
    }  
}
```





# Processor API (PAPI) Example - II

@Override

```
public void process(Record<String, String> record) {  
    final String[] words = record.value().toLowerCase().split("\\W+");  
  
    for (final String word : words) {  
        Long oldVal = kvStore.get(word);  
        if (oldVal == null) {  
            oldVal = 0L;  
        }  
        kvStore.put(word, oldVal + 1);  
        context.forward(new Record<>(  
            word,  
            String.format("%-15s -> %4d", word, oldVal + 1),  
            record.timestamp()  
        ));  
    }  
}
```





# Streams DSL: Stateless Transformations

Stateless transformations do not require state for processing and they do not require a state store associated with the stream processor. Kafka allows you to materialize the result from a stateless KTable transformation. This allows the result to be queried through interactive queries. To materialize a KTable, each of stateless operations can be augmented with an optional queryableStoreName argument:

**Branch:** KStream → BranchedKStream

**Filter:** KStream → Kstream, **Filter:** KTable → Ktable

**Inverse Filter filterNot:** KStream → Kstream, **filterNot:** KTable → Ktable

**FlatMap:** KStream → Kstream, **FlatMap (values only):** KStream → Kstream

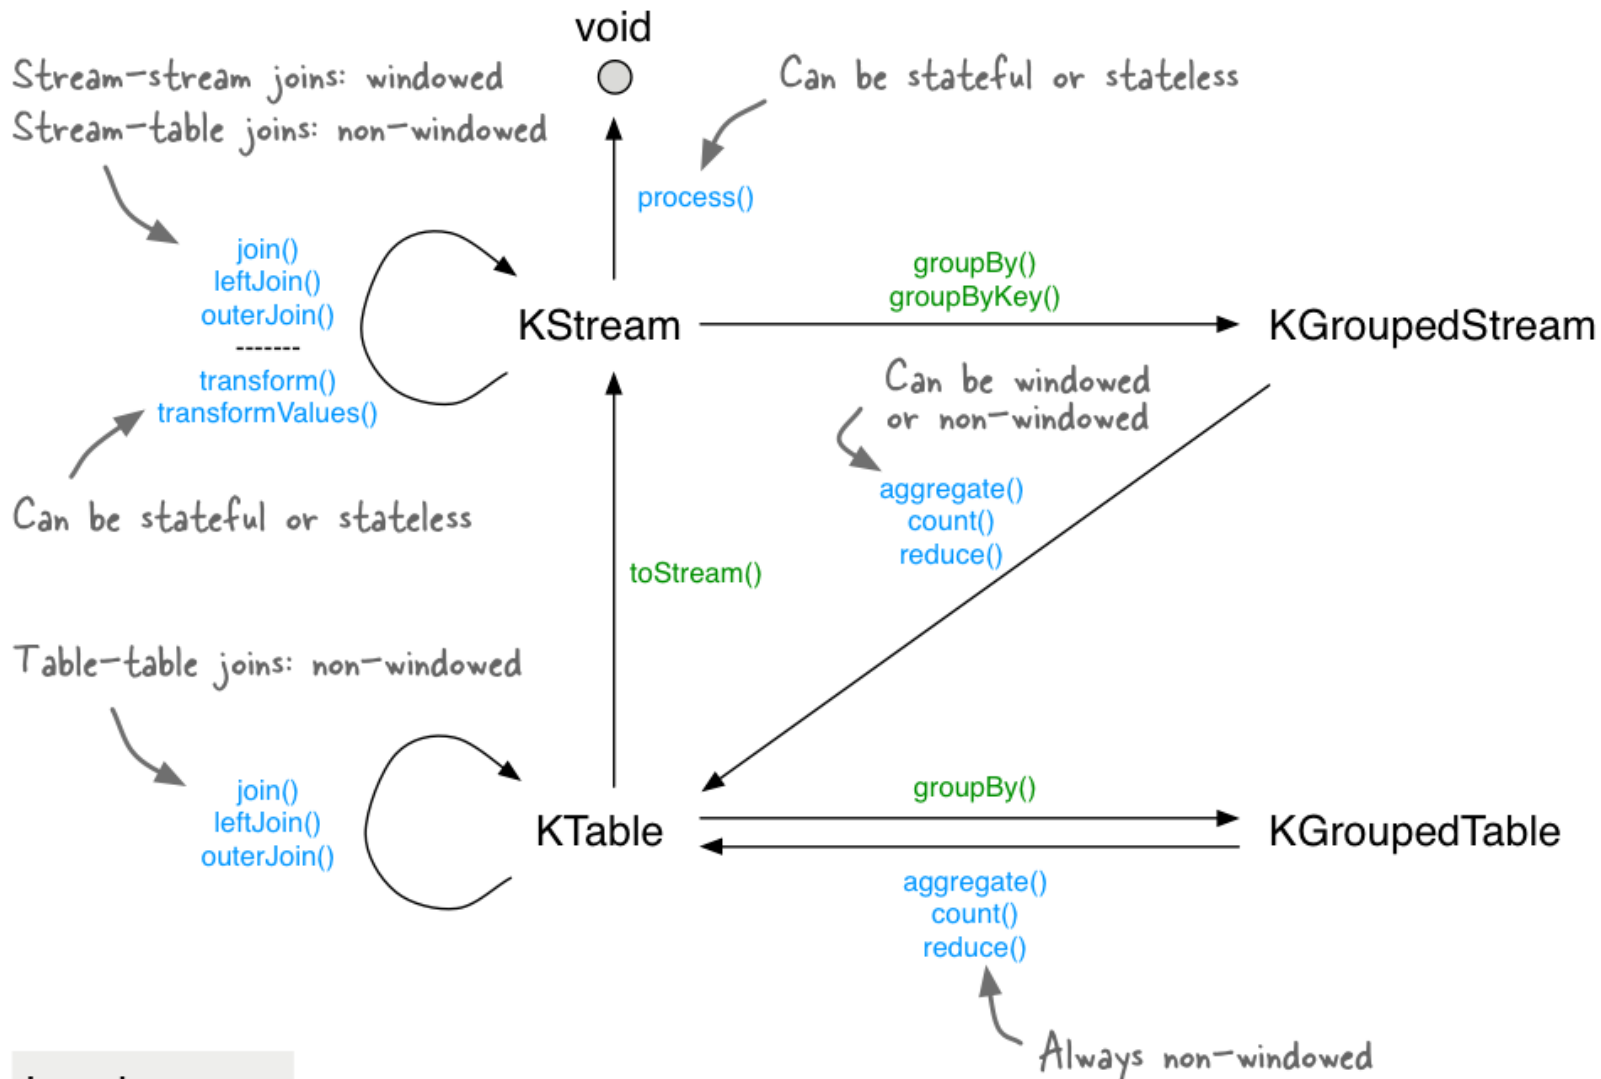
**Foreach:** KStream → void | KStream → void | KTable → void

**GroupByKey:** KStream → KGroupedStream, **GroupBy:** KStream → KGroupedStream





# Streams DSL: Stateful Transformations



## Legend

Stateful operations  
Stateless operations

GlobalKTable  
no direct operations



# Streams DSL: Stateful Transformations

- **Stateful transformations** depend on **state** for processing inputs and producing outputs and **require a state store** associated with the stream processor. In aggregating operations, a **windowing state store** is used to collect the **latest aggregation results per window**. In join operations, a **windowing state store** is used to collect all of records received **within the defined window boundary**.
- **non-windowed** aggregations and non-windowed **KTables** use **TimestampedKeyValueStores**
- **time-windowed** aggregations and **KStream-KStream joins** use **TimestampedWindowStores**
- **session windowed** aggregations use **SessionStores** (there is no timestamped session store as of now)
- **State stores** are **fault-tolerant**. In case of failure, Kafka Streams guarantees to fully restore all state stores prior to resuming the processing.





# Streams DSL: Stateful Transformations

- **Stateful transformations** depend on **state** for processing inputs and producing outputs and **require a state store** associated with the stream processor. In aggregating operations, a **windowing state store** is used to collect the **latest aggregation results per window**. In join operations, a **windowing state store** is used to collect all of records received **within the defined window boundary**.
- **non-windowed** aggregations and non-windowed **KTables** use **TimestampedKeyValueStores**
- **time-windowed** aggregations and **KStream-KStream joins** use **TimestampedWindowStores**
- **session windowed** aggregations use **SessionStores** (there is no timestamped session store as of now)
- **State stores** are **fault-tolerant**. In case of failure, Kafka Streams guarantees to fully restore all state stores prior to resuming the processing.





# Types of Stateful Transformations

Available stateful transformations in the DSL include:

- Aggregating
- Joining
- Windowing (as part of aggregations and joins)
- Applying custom processors and transformers, which may be stateful, for Processor API integration



# Aggregating

- After records are grouped by key via `groupByKey` or `groupByKey` – and thus represented as either a `KGroupedStream` or a `KGroupedTable`, they can be **aggregated** via an operation such as `reduce`. Aggregations are key-based operations, which means that they **always operate over records** (notably record values) of the **same key**. You can perform aggregations on windowed or non-windowed data.
- Types of windows:

Window name	Behavior	Short description
<a href="#">Hopping time window</a>	Time-based	Fixed-size, overlapping windows
<a href="#">Tumbling time window</a>	Time-based	Fixed-size, non-overlapping, gap-less windows
<a href="#">Sliding time window</a>	Time-based	Fixed-size, overlapping windows that work on differences between record timestamps
<a href="#">Session window</a>	Session-based	Dynamically-sized, non-overlapping, data-driven windows





# Hopping Windows

```
import java.time.Duration;  
import org.apache.kafka.streams.kstream.TimeWindows;
```


```
// A hopping time window with a size of 5 minutes and an advance interval of 1 min.  
// The window's name -- the string parameter -- is used to e.g. name the backing state store.  
Duration windowSize = Duration.ofMinutes(5);  
Duration advance = Duration.ofMinutes(1);  
TimeWindows.ofSizeWithNoGrace(windowSize).advanceBy(advance);
```







## A 5-min Hopping Window with a 1-min "hop"





# Tumbling Time Windows

```
import java.time.Duration;
import org.apache.kafka.streams.kstream.TimeWindows;



// A tumbling time window with a size of 5 minutes (and, by definition, an implicit
// advance interval of 5 minutes), and grace period of 1 minute.
Duration windowSize = Duration.ofMinutes(5);
Duration gracePeriod = Duration.ofMinutes(1);
TimeWindows.ofSizeAndGrace(windowSize, gracePeriod);

// The above is equivalent to the following code:
TimeWindows.ofSizeAndGrace(windowSize, gracePeriod).advanceBy(windowSize);
```



# A 5-min Tumbling Window

Data records



Windows are created per record key





# Sliding Time Windows

```
import org.apache.kafka.streams.kstream.SlidingWindows;
```

```
// A sliding time window with a time difference of 10 minutes and grace period of 30 minutes
```

```
Duration timeDifference = Duration.ofMinutes(10);
```

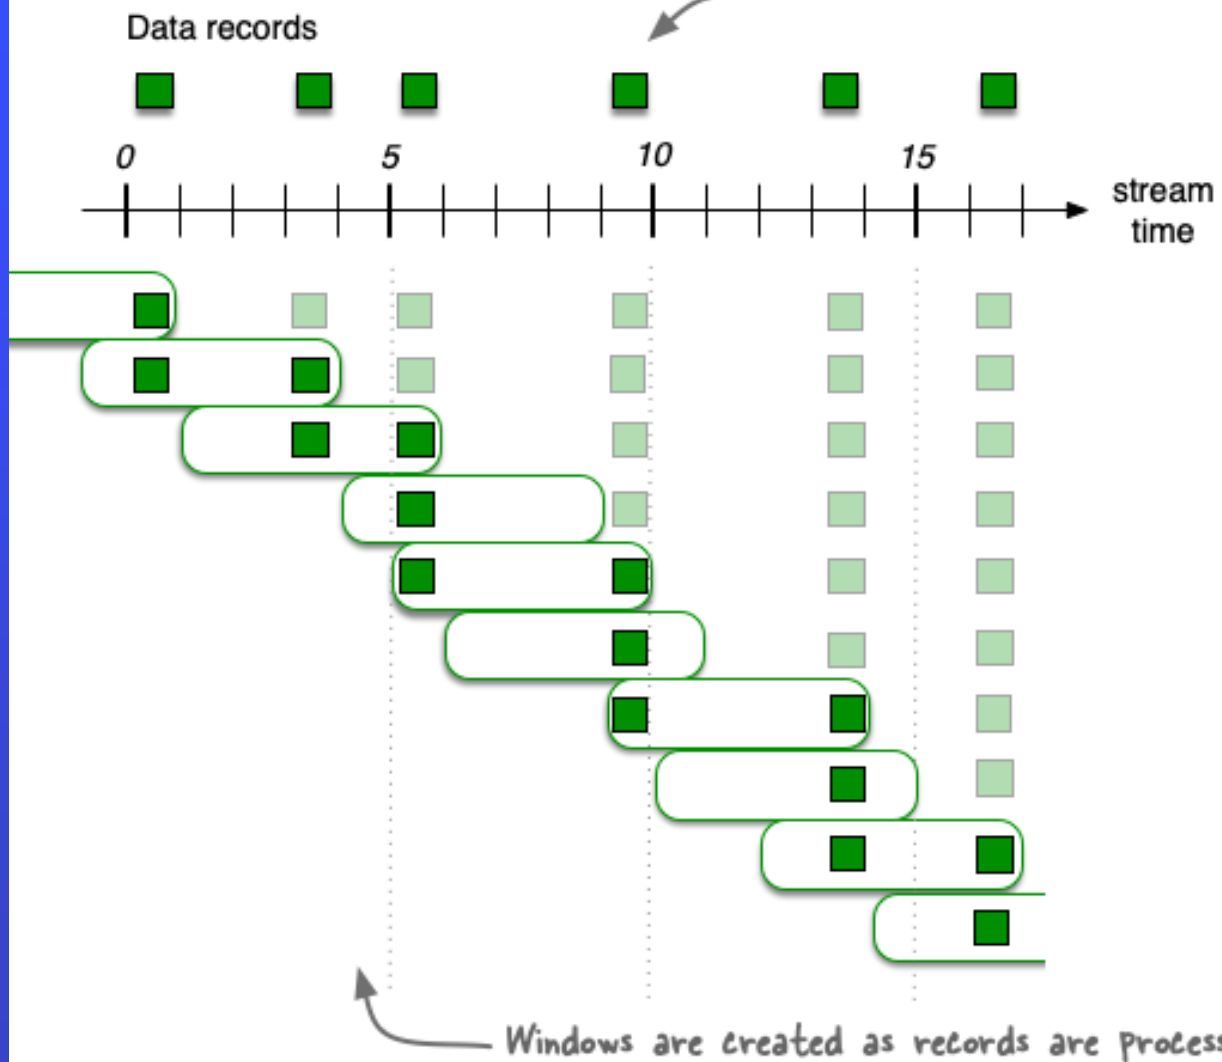
```
Duration gracePeriod = Duration.ofMinutes(30);
```

```
SlidingWindows.ofTimeDifferenceAndGrace(timeDifference, gracePeriod);
```



## A 5-ms Sliding Window

Same colour means same record key





# Session Windows

```
import java.time.Duration;  
import org.apache.kafka.streams.kstream.SessionWindows;  
  
// A session window with an inactivity gap of 5 minutes.  
SessionWindows.ofInactivityGapWithNoGrace(Duration.ofMinutes(5));
```




## A Session Window with a 5-min inactivity gap

Data records



Same colour means same record key

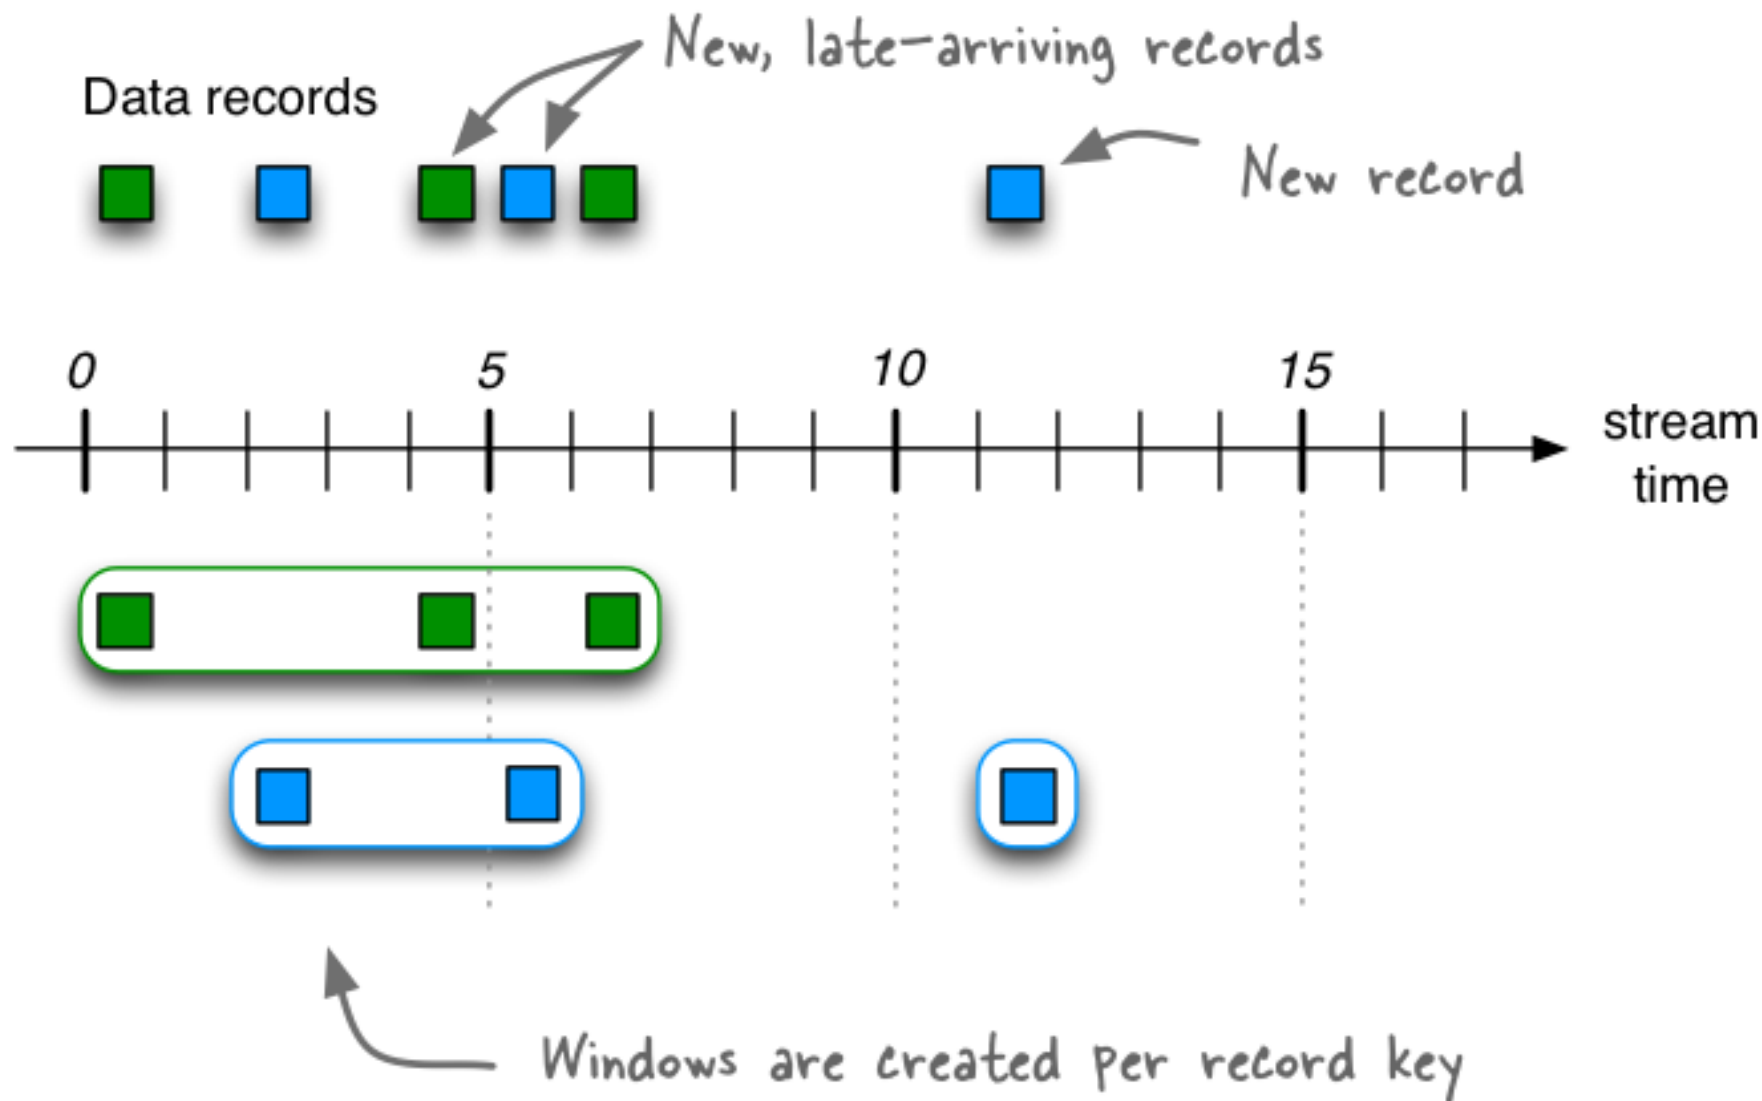


Windows are created per record key





## A Session Window with a 5-min inactivity gap





# Window Final Results

- In Kafka Streams, windowed computations **update their results continuously**. As new data arrives for a window, freshly computed results are emitted downstream.
- However, some applications need to take action only on the final result of a windowed computation. Common examples of this are sending alerts or delivering results to a system that doesn't support updates.

```
KGroupedStream<UserId, Event> grouped = ...;  
grouped  
    .windowedBy(TimeWindows.ofSizeAndGrace(Duration.ofHours(1), Duration.ofMinutes(10)))  
    .count()  
    .suppress(Suppressed.untilWindowCloses(unbounded()))  
    .filter((windowedUserId, count) -> count < 3)  
    .toStream()  
    .foreach((windowedUserId, count) -> sendAlert(windowedUserId.window(), windowedUserId.key(), count));
```





# Controlling KTable Emit Rate

- some applications need to take other actions, such as calling out to external systems, and therefore need to exercise some control over the rate of invocations, for example of **KStream#foreach**.
- Rather than achieving this as a **side-effect of the KTable record cache**, you can directly impose a rate limit via the **KTable#suppress** operator.

```
KGroupedTable<String, String> groupedTable = ...;  
groupedTable  
    .count()  
    .suppress(untilTimeLimit(ofMinutes(5), maxBytes(1_000_000L).emitEarlyWhenFull()))  
    .toStream()  
    .foreach((key, count) -> updateCountsDatabase(key, count));
```





# Joining

Join operands	Type	(INNER) JOIN	LEFT JOIN	OUTER JOIN
KStream-to-KStream	Windowed	Supported	Supported	Supported
KTable-to-KTable	Non-windowed	Supported	Supported	Supported
KTable-to-KTable Foreign-Key Join	Non-windowed	Supported	Supported	Not Supported
KStream-to-KTable	Non-windowed	Supported	Supported	Not Supported
KStream-to-GlobaKTable	Non-windowed	Supported	Supported	Not Supported
KTable-to-GlobaKTable	N/A	Not Supported	Not Supported	Not Supported



# Join Co-partitioning Requirements

- For **equi-joins**, input data must be co-partitioned when joining. This ensures that input records with the same key from both sides of the join, are delivered to the same stream task during processing.
- Co-partitioning is not required when performing **KTable-KTable Foreign-Key joins** and **Global KTable joins**.
- The input topics of the join (left side and right side) must have the **same number of partitions**.
- All applications that write to the input topics must have the **same partitioning strategy** so that records with the same key are delivered to same partition number. In other words, the keyspace of the input data must be distributed across partitions in the same manner.





# Join Co-partitioning Requirements - II


- Why is data co-partitioning required? Because **KStream-KStream**, **KTable-KTable**, and **KStream-KTable** joins are performed based on the keys of records (e.g., `leftRecord.key == rightRecord.key`), it is required that the input streams/tables of a join are **co-partitioned by key**.
- There are two exceptions where co-partitioning is not required. For **KStream-GlobalKTable** joins, co-partitioning is not required because all partitions of the **GlobalKTable**'s underlying changelog stream are made available to each **KafkaStreams** instance. That is, each instance has a **full copy** of the changelog stream. Further, a **KeyValueMapper** allows for non-key based joins from the **KStream** to the **GlobalKTable**. **KTable-KTable Foreign-Key** joins also do not require co-partitioning. Kafka Streams internally ensures co-partitioning for **Foreign-Key** joins.



# Demos

Available @ Github:

<https://github.com/iproduct/kafka-streams-devbg>



Следете актуалните обяви за Java

DEV.BG



# Thank you!

Contacts:


 <https://www.linkedin.com/in/trayaniliev/>

 <https://www.facebook.com/IPT.EACAD>

 <https://github.com/iproduct>

 <https://twitter.com/trayaniliev>

СЛЕДВАЩО СЪБИТИЕ



<https://dev.bg/groups/java/>

Лектор

Дата

Език



Следете актуалните обяви за **Java**

**DEV.BG**